

DEEP NEURAL NETWORKS PROJECT

Rama Alyoubi 2110112
Rimas Alshehri 2110240

PROJECT INTRODUCTION

Objective :

Develop a neural network tailored for the precise recognition of individual handwritten Arabic characters.

Challenge :

Our task is to engineer and fine-tune a deep neural network architecture that excels in accuracy when confronted with the test set.

Journey :

- From problem introduction to neural network intricacies.
- Through dataset details to results and conclusions.

DATA OVERVIEW

اپنے شجاع خود رکھ لیں،
خوب سرگشہ و لایں،

اب تھاتھ دذراز طکل من صضع

01 - DATA DESCRIPTION

The Arabic Chars MNIST dataset is a collection of 28x28 grayscale images of handwritten Arabic characters, each labeled for classification.

02 - NATURE OF THE DATA

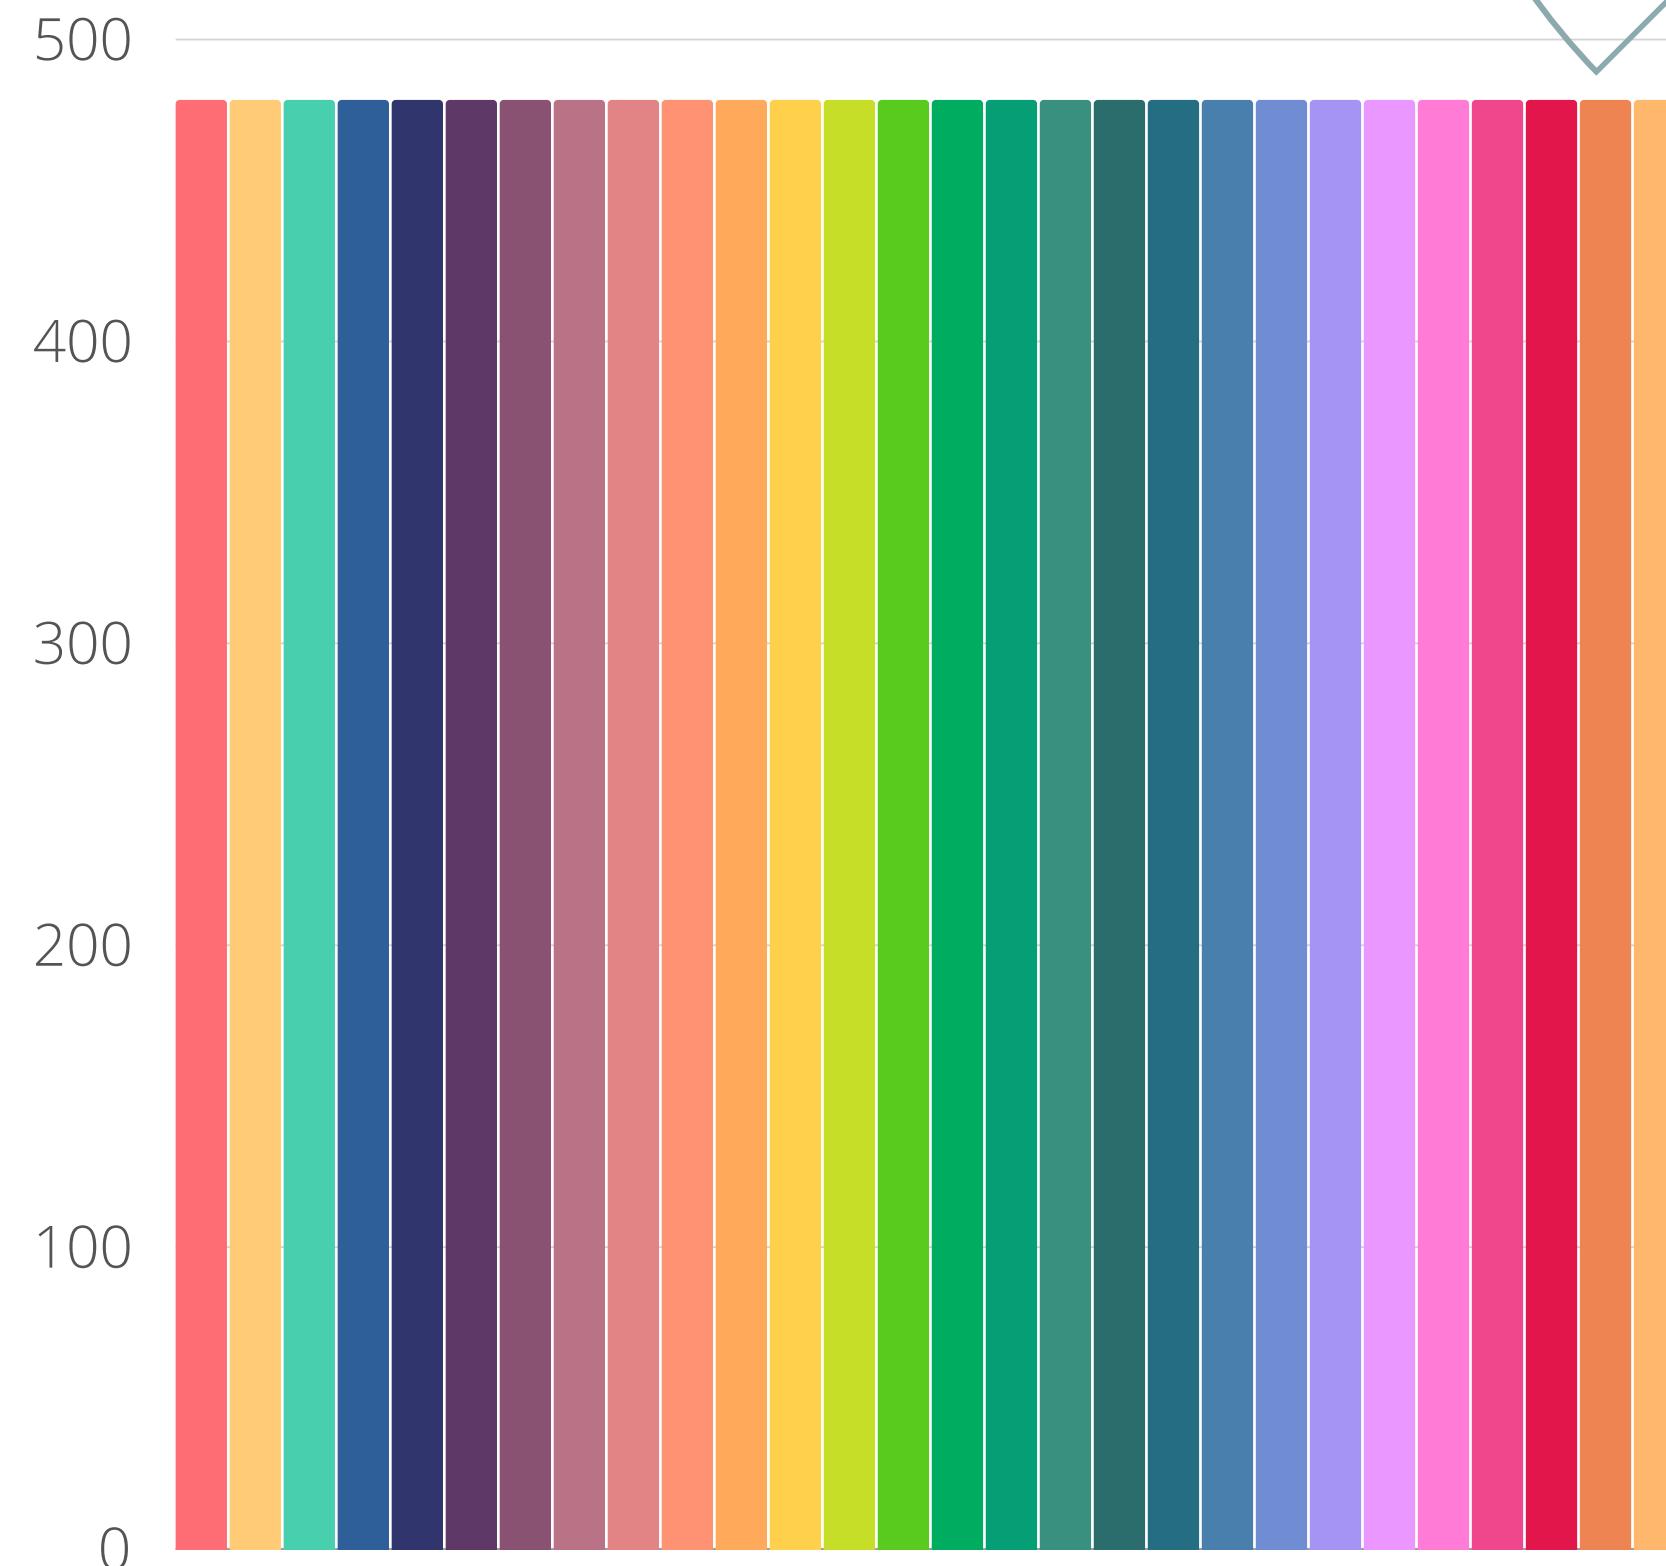
Grayscale images in a standardized size, featuring a range of handwriting styles, suitable for training OCR algorithms.

03 - SIZE OF THE DATA

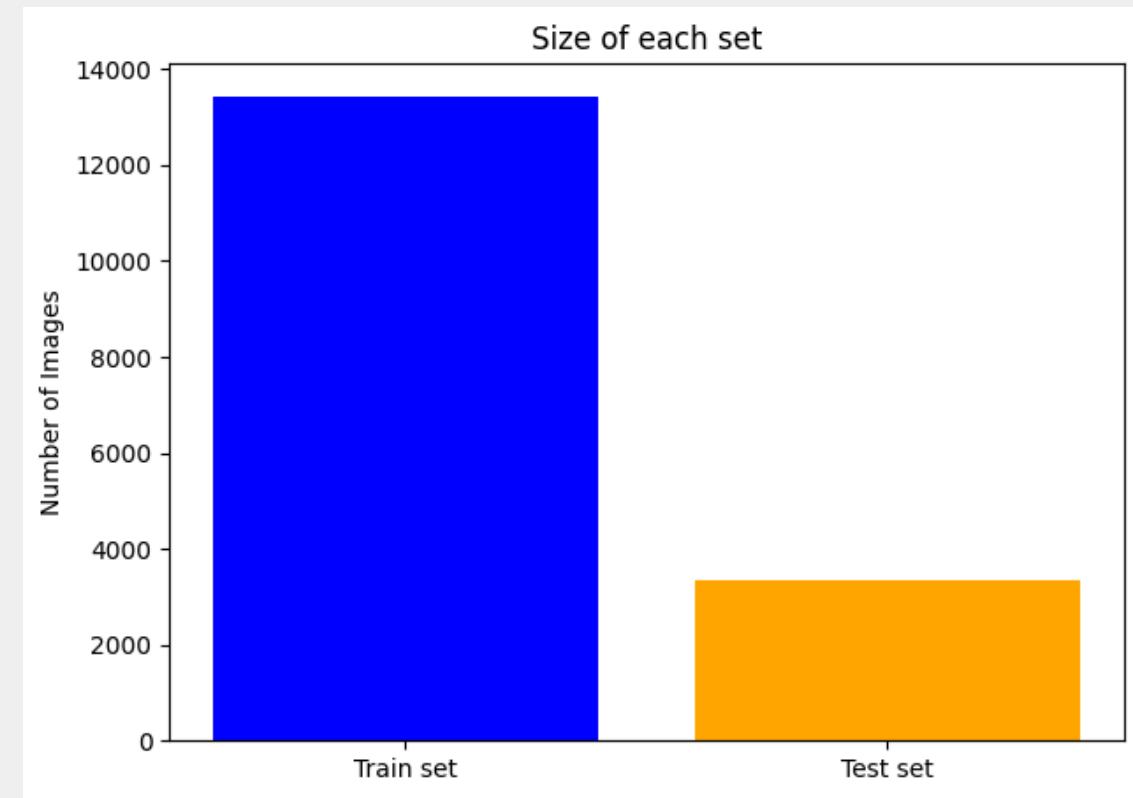
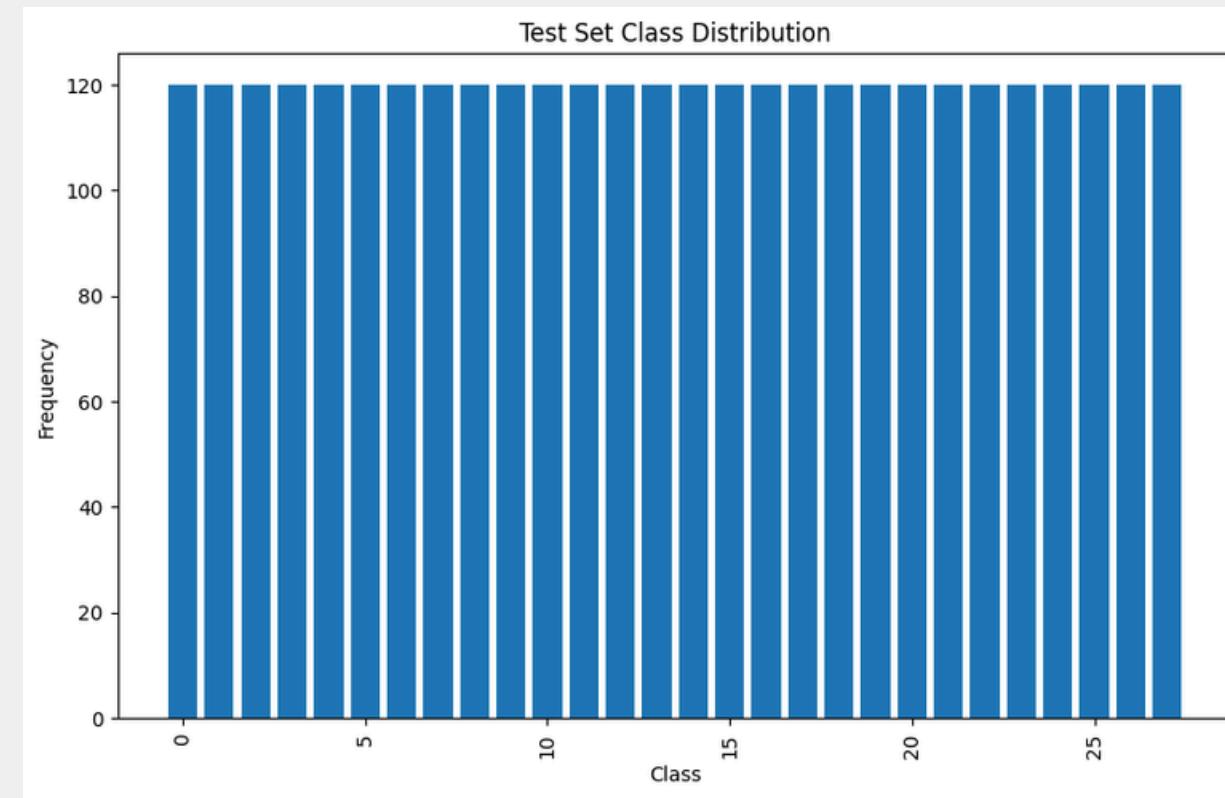
Comprising 16,800 images in total, with 13,440 for training and 3,360 for testing across 28 classes.

EXPLORATORY DATA ANALYSIS (EDA)

In our comprehensive EDA, we've confirmed that the dataset is perfectly balanced, with each of the 28 Arabic character classes represented by 480 images. This uniform distribution is ideal as it prevents model bias towards more frequently represented classes and ensures equal learning opportunity for all characters.

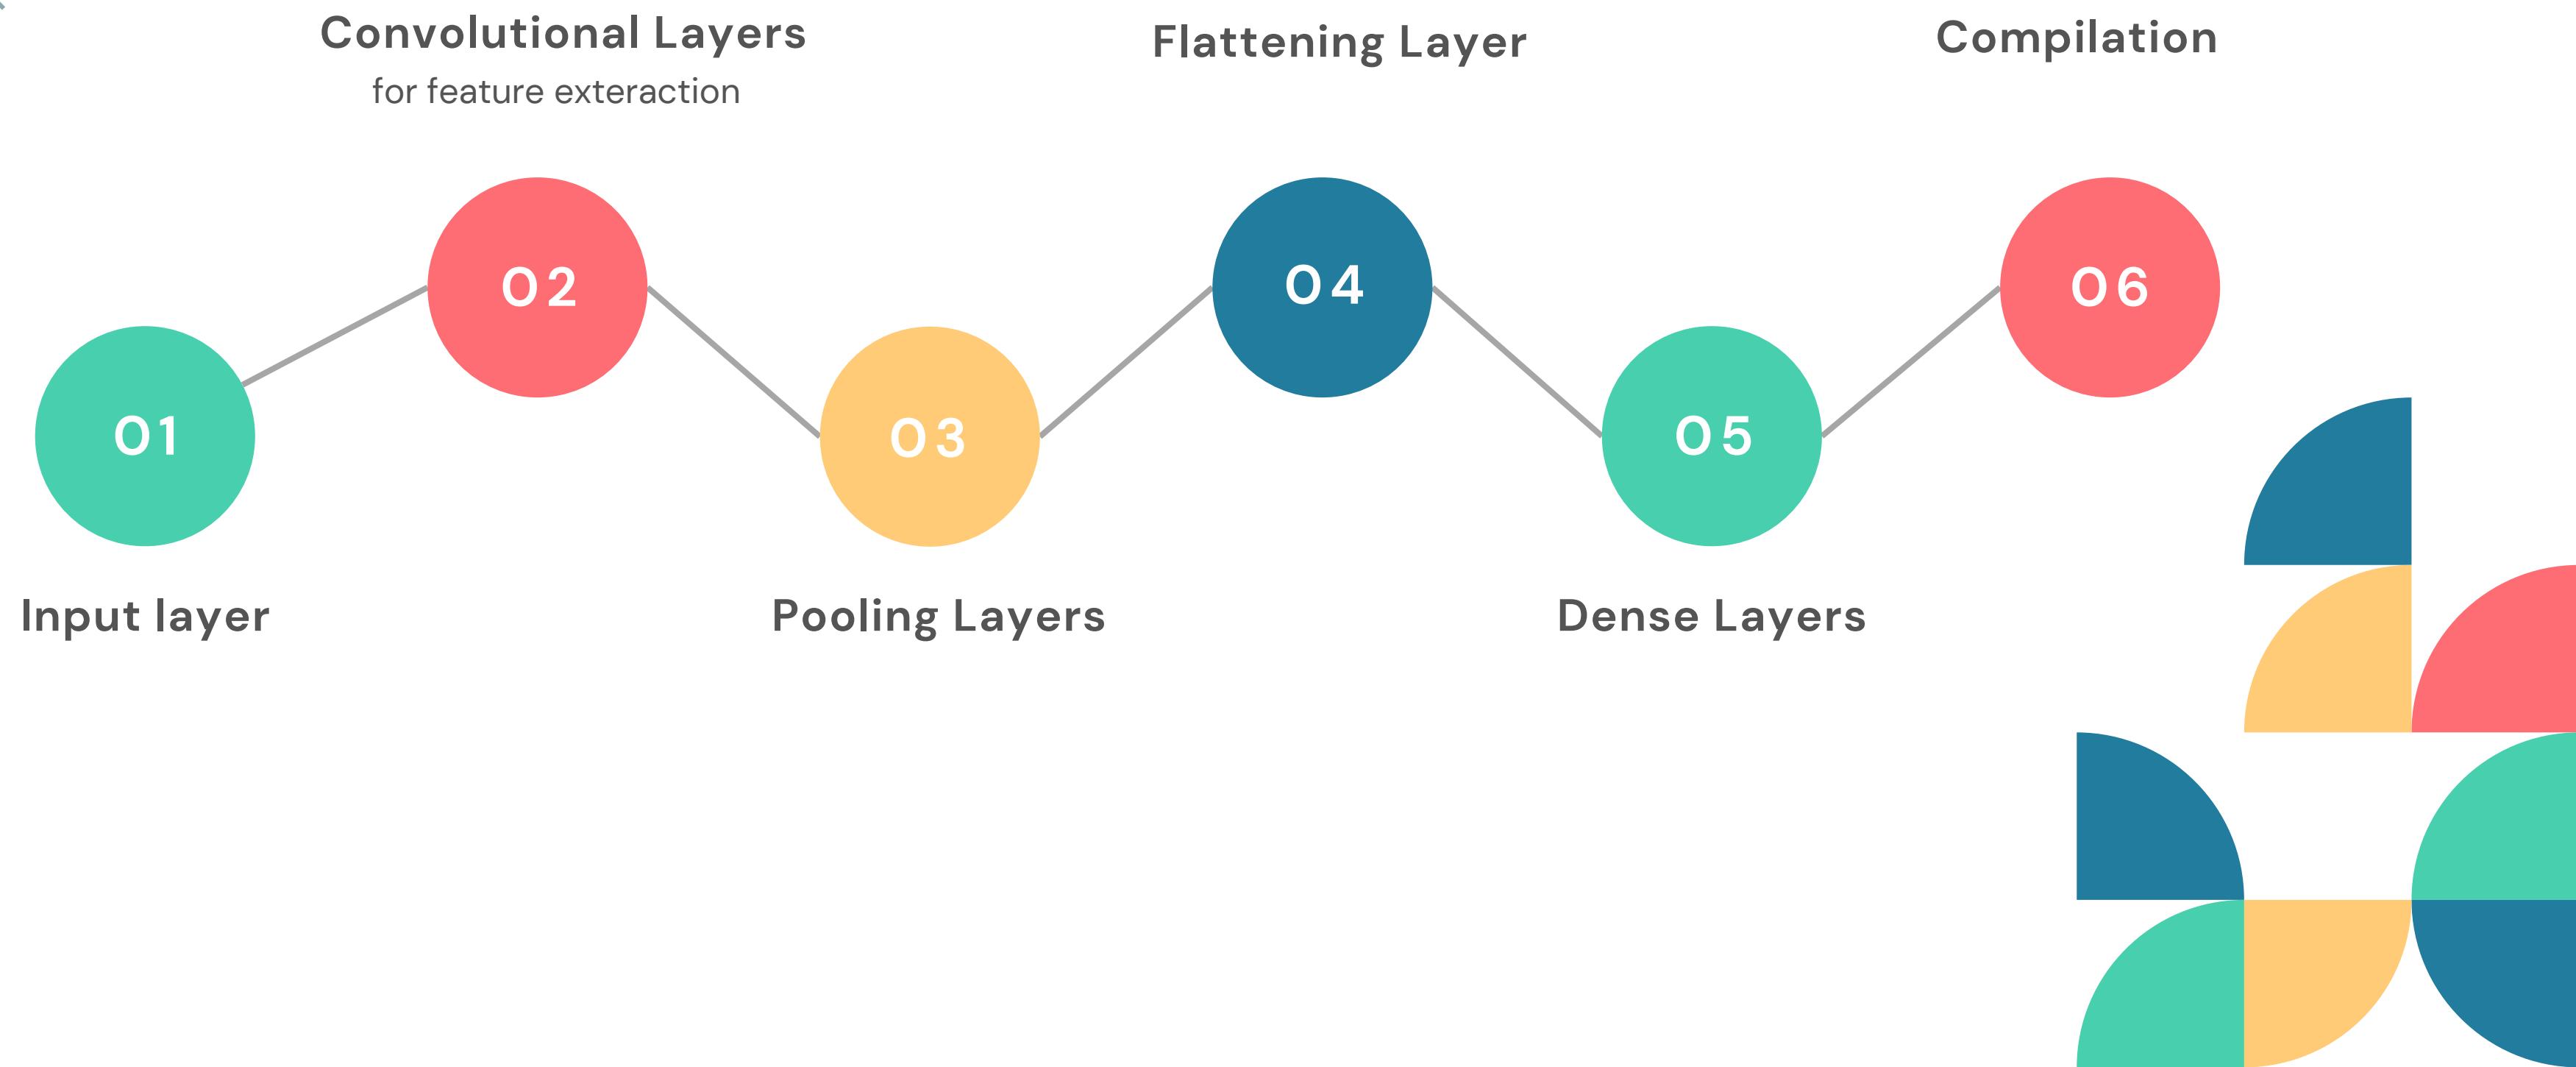


DATA SPLITTING & VIZUALIZATION

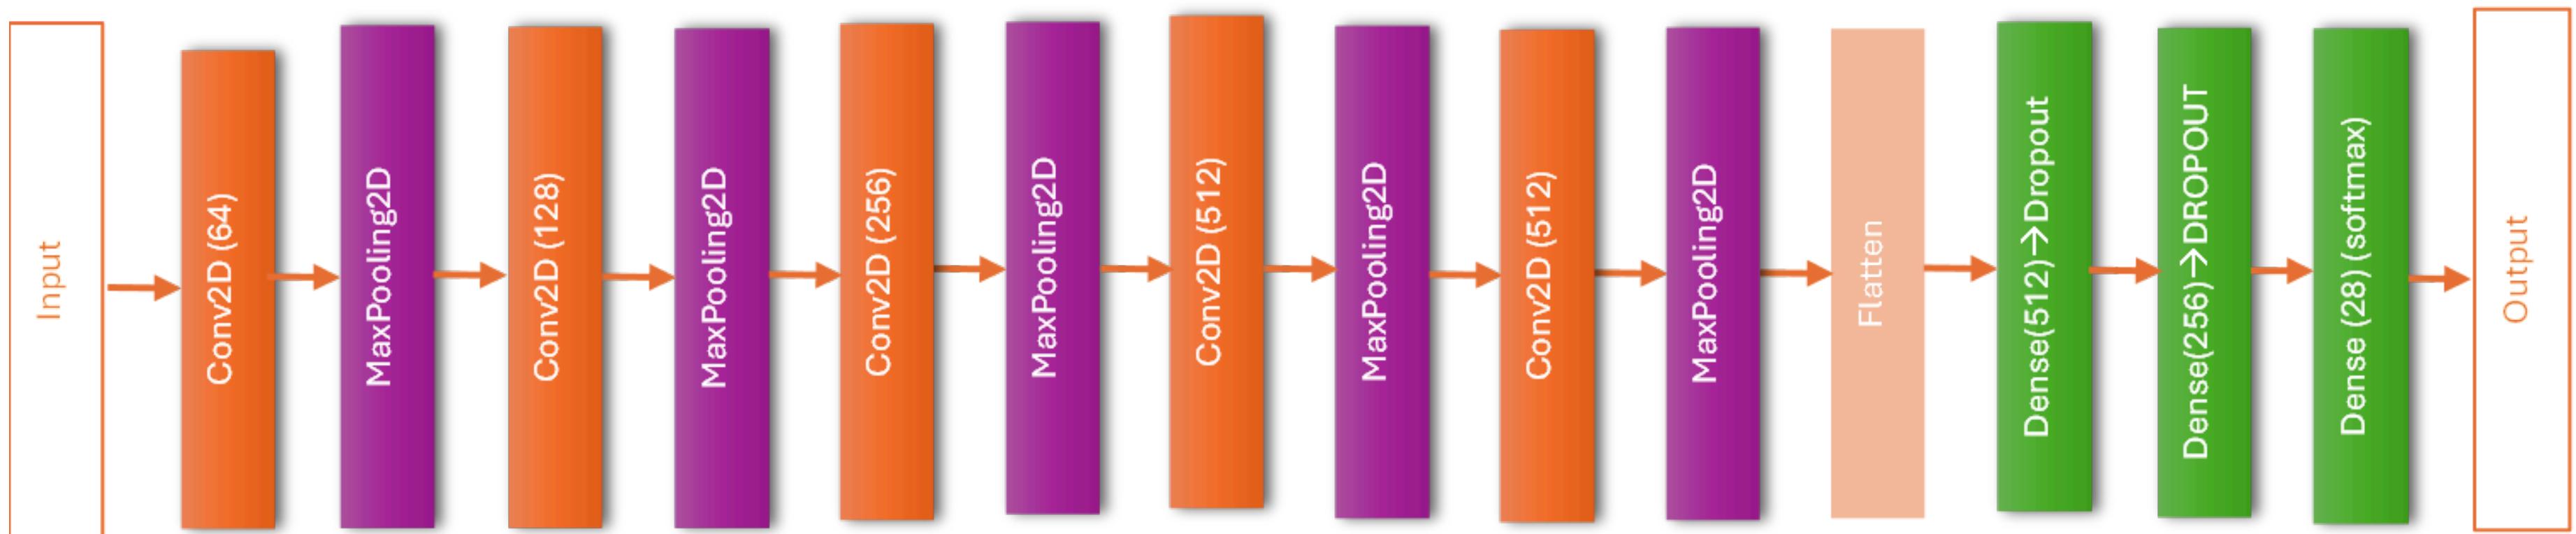


PROCESS OF BUILDING THE CNN MODEL

BASIC STRUCTURE



MODEL ARCHITECTURE



IMPLEMENTATION

```
def create_advanced_cnn_model(input_shape, num_classes):
    model = Sequential([
        # First conv layer
        Conv2D(64, kernel_size=(3, 3), strides=(1, 1), activation='relu', input_shape=input_shape, padding='same', kernel_regularizer=l2(0.001)),
        MaxPooling2D(pool_size=(2, 2)),

        # Second conv layer
        Conv2D(128, kernel_size=(3, 3), strides=(1, 1), activation='relu', padding='same', kernel_regularizer=l2(0.001)),
        MaxPooling2D(pool_size=(2, 2)),

        # Third conv layer
        Conv2D(256, kernel_size=(3, 3), strides=(1, 1), activation='relu', padding='same', kernel_regularizer=l2(0.001)),
        MaxPooling2D(pool_size=(2, 2)),

        # Fourth conv layer
        Conv2D(512, kernel_size=(3, 3), strides=(1, 1), activation='relu', padding='same', kernel_regularizer=l2(0.001)),
        MaxPooling2D(pool_size=(2, 2)),

        # Fifth conv layer
        Conv2D(512, kernel_size=(3, 3), strides=(1, 1), activation='relu', padding='same', kernel_regularizer=l2(0.001)),
        MaxPooling2D(pool_size=(2, 2)),

        # Flattening the layers
        Flatten(),

        # Dense layers
        Dense(512, activation='relu', kernel_regularizer=l2(0.001)),
        Dropout(0.5),
        Dense(256, activation='relu', kernel_regularizer=l2(0.001)),
        Dropout(0.5),
        Dense(num_classes, activation='softmax')
    ])

    model.compile(optimizer=Adam(), loss='categorical_crossentropy', metrics=['accuracy'])
    return model
```

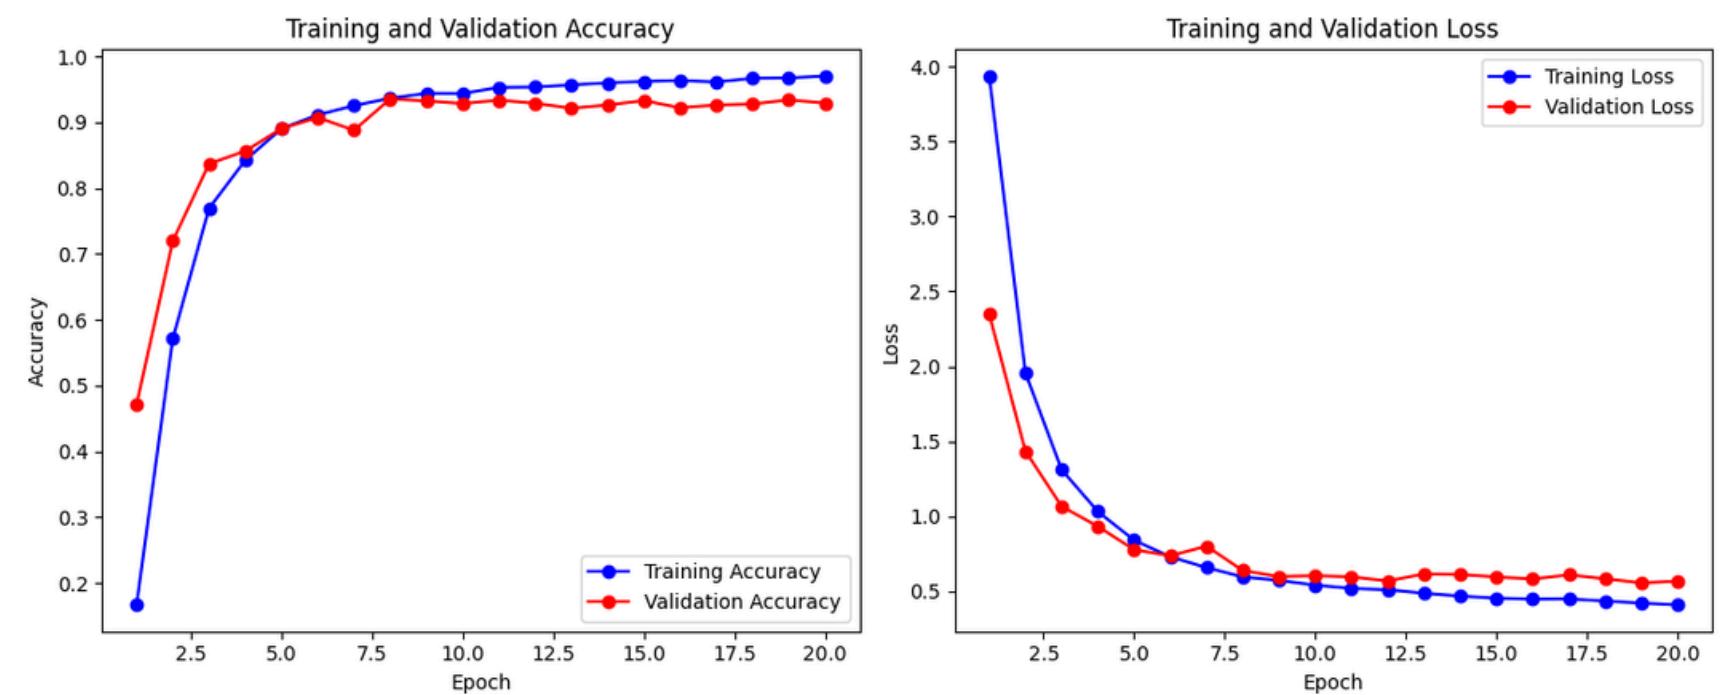
MODEL SUMMARY

Model: "sequential"		
Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 64)	640
max_pooling2d (MaxPooling2D)	(None, 16, 16, 64)	0
conv2d_1 (Conv2D)	(None, 16, 16, 128)	73856
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 128)	0
conv2d_2 (Conv2D)	(None, 8, 8, 256)	295168
max_pooling2d_2 (MaxPooling2D)	(None, 4, 4, 256)	0
conv2d_3 (Conv2D)	(None, 4, 4, 512)	1180160
max_pooling2d_3 (MaxPooling2D)	(None, 2, 2, 512)	0
conv2d_4 (Conv2D)	(None, 2, 2, 512)	2359808
max_pooling2d_4 (MaxPooling2D)	(None, 1, 1, 512)	0
flatten (Flatten)	(None, 512)	0
dense (Dense)	(None, 512)	262656
dropout (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 256)	131328
dropout_1 (Dropout)	(None, 256)	0
dense_2 (Dense)	(None, 28)	7196
<hr/>		
Total params: 4310812 (16.44 MB)		
Trainable params: 4310812 (16.44 MB)		
Non-trainable params: 0 (0.00 Byte)		

TRAINING VS VALIDATION

```
history = model.fit(  
    train_images, train_labels_one_hot,  
    epochs=20,  
    batch_size=32,  
    validation_split=0.2, # Using 20% of the data for validation  
    verbose=1  
)  
  
Epoch 1/20  
336/336 [=====] - 12s 11ms/step - loss: 3.9375 - accuracy: 0.1667 - val_loss: 2.3484 - val_accuracy: 0.4714  
Epoch 2/20  
336/336 [=====] - 4s 11ms/step - loss: 1.9556 - accuracy: 0.5714 - val_loss: 1.4337 - val_accuracy: 0.7199  
Epoch 3/20  
336/336 [=====] - 4s 12ms/step - loss: 1.3098 - accuracy: 0.7693 - val_loss: 1.0653 - val_accuracy: 0.8371  
Epoch 4/20  
336/336 [=====] - 4s 12ms/step - loss: 1.0307 - accuracy: 0.8424 - val_loss: 0.9318 - val_accuracy: 0.8564  
Epoch 5/20  
336/336 [=====] - 4s 11ms/step - loss: 0.8395 - accuracy: 0.8905 - val_loss: 0.7761 - val_accuracy: 0.8906  
Epoch 6/20  
336/336 [=====] - 3s 10ms/step - loss: 0.7288 - accuracy: 0.9114 - val_loss: 0.7363 - val_accuracy: 0.9074  
Epoch 7/20  
336/336 [=====] - 4s 12ms/step - loss: 0.6567 - accuracy: 0.9254 - val_loss: 0.7999 - val_accuracy: 0.8880  
Epoch 8/20  
336/336 [=====] - 4s 12ms/step - loss: 0.5949 - accuracy: 0.9367 - val_loss: 0.6371 - val_accuracy: 0.9356  
Epoch 9/20  
336/336 [=====] - 4s 11ms/step - loss: 0.5715 - accuracy: 0.9441 - val_loss: 0.5971 - val_accuracy: 0.9327  
Epoch 10/20  
336/336 [=====] - 4s 11ms/step - loss: 0.5399 - accuracy: 0.9439 - val_loss: 0.6029 - val_accuracy: 0.9286  
Epoch 11/20  
336/336 [=====] - 4s 12ms/step - loss: 0.5180 - accuracy: 0.9531 - val_loss: 0.5952 - val_accuracy: 0.9338  
Epoch 12/20  
336/336 [=====] - 4s 12ms/step - loss: 0.5090 - accuracy: 0.9537 - val_loss: 0.5677 - val_accuracy: 0.9289  
Epoch 13/20  
336/336 [=====] - 4s 10ms/step - loss: 0.4858 - accuracy: 0.9569 - val_loss: 0.6146 - val_accuracy: 0.9215  
Epoch 14/20  
336/336 [=====] - 4s 11ms/step - loss: 0.4669 - accuracy: 0.9600 - val_loss: 0.6110 - val_accuracy: 0.9263  
Epoch 15/20  
336/336 [=====] - 4s 11ms/step - loss: 0.4530 - accuracy: 0.9622 - val_loss: 0.5946 - val_accuracy: 0.9330  
Epoch 16/20  
336/336 [=====] - 4s 12ms/step - loss: 0.4464 - accuracy: 0.9636 - val_loss: 0.5811 - val_accuracy: 0.9222  
Epoch 17/20  
336/336 [=====] - 4s 11ms/step - loss: 0.4480 - accuracy: 0.9614 - val_loss: 0.6084 - val_accuracy: 0.9263  
Epoch 18/20  
336/336 [=====] - 4s 11ms/step - loss: 0.4340 - accuracy: 0.9670 - val_loss: 0.5824 - val_accuracy: 0.9278  
Epoch 19/20  
336/336 [=====] - 4s 12ms/step - loss: 0.4200 - accuracy: 0.9674 - val_loss: 0.5550 - val_accuracy: 0.9342  
Epoch 20/20  
336/336 [=====] - 4s 12ms/step - loss: 0.4081 - accuracy: 0.9705 - val_loss: 0.5666 - val_accuracy: 0.9293
```

training accuracy: 0.97
validation accuracy: 0.92

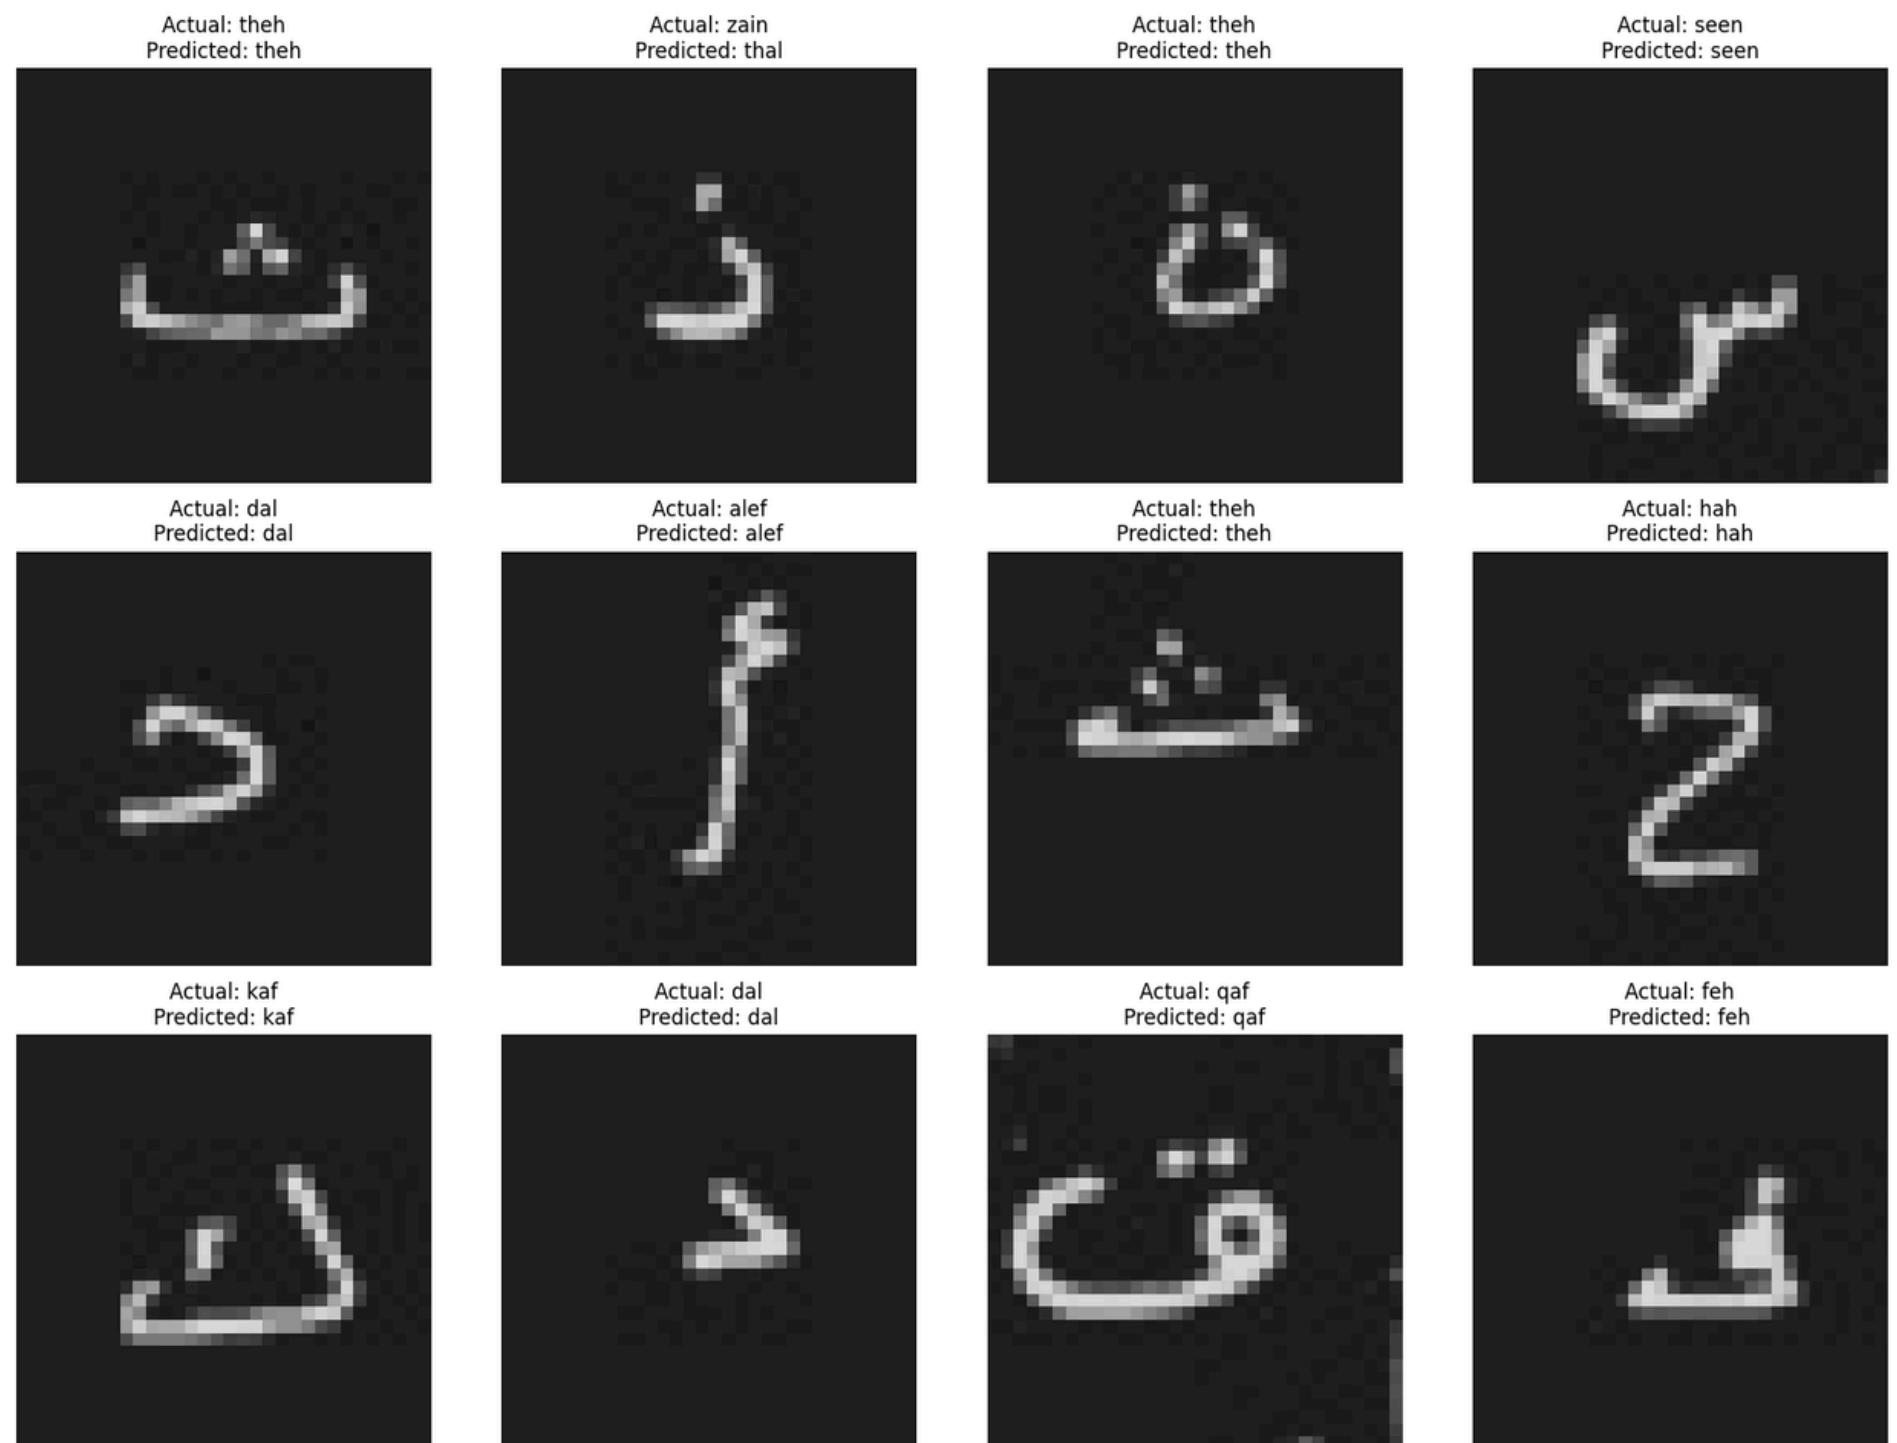


TESTING

```
# Evaluate the model on the test set
test_loss, test_accuracy = model.evaluate(test_images, test_labels_one_hot)
print(f"Test Loss: {test_loss}")
print(f"Test Accuracy: {test_accuracy}")

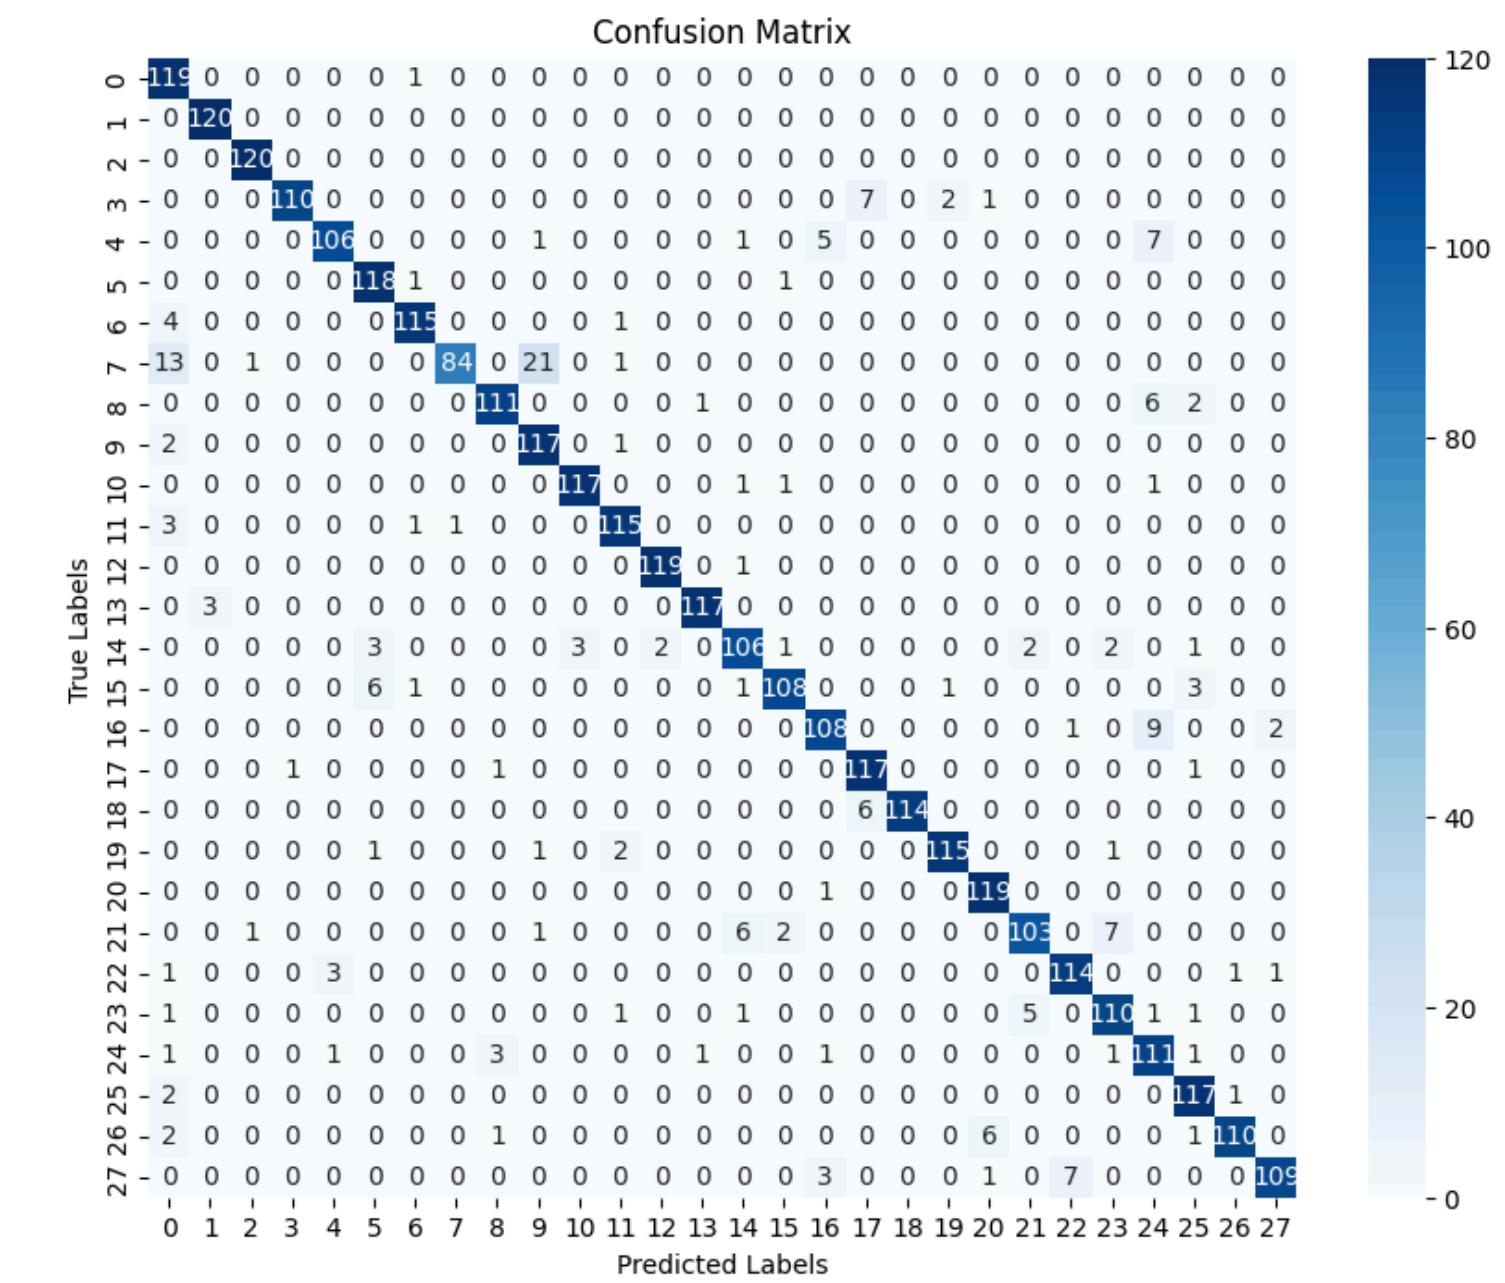
105/105 [=====] - 0s 4ms/step - loss: 0.5559 - accuracy: 0.9372
Test Loss: 0.5559284090995789
Test Accuracy: 0.9372023940086365
```

testing accuracy: 0.93



CLASSIFICATION REPORT & CONFUSION MATRIX

105/105 [=====] - 0s 3ms/step				
	precision	recall	f1-score	support
0	0.80	0.99	0.89	120
1	0.98	1.00	0.99	120
2	0.98	1.00	0.99	120
3	0.99	0.92	0.95	120
4	0.96	0.88	0.92	120
5	0.92	0.98	0.95	120
6	0.97	0.96	0.96	120
7	0.99	0.70	0.82	120
8	0.96	0.93	0.94	120
9	0.83	0.97	0.90	120
10	0.97	0.97	0.97	120
11	0.95	0.96	0.95	120
12	0.98	0.99	0.99	120
13	0.98	0.97	0.98	120
14	0.91	0.88	0.89	120
15	0.96	0.90	0.93	120
16	0.92	0.90	0.91	120
17	0.90	0.97	0.94	120
18	1.00	0.95	0.97	120
19	0.97	0.96	0.97	120
20	0.94	0.99	0.96	120
21	0.94	0.86	0.90	120
22	0.93	0.95	0.94	120
23	0.91	0.92	0.91	120
24	0.82	0.93	0.87	120
25	0.92	0.97	0.95	120
26	0.98	0.92	0.95	120
27	0.97	0.91	0.94	120
accuracy			0.94	3360
macro avg	0.94	0.94	0.94	3360
weighted avg	0.94	0.94	0.94	3360



DIFFERENT VARIATIONS OF PARAMETERS TO THE MODEL.

CHANGE 1 TO THE MODEL

changing convolution size, kernel size, L2 regularization

IMPLEMENTATION

```
def create_advanced_cnn_model(input_shape, num_classes):
    model = Sequential([
        # First conv layer
        Conv2D(64, kernel_size=(3, 3), strides=(1, 1), activation='relu', input_shape=input_shape, padding='same', kernel_regularizer=l2(0.002)),
        MaxPooling2D(pool_size=(2, 2)),

        # Second conv layer
        Conv2D(32, kernel_size=(6, 6), strides=(1, 1), activation='relu', padding='same', kernel_regularizer=l2(0.002)),
        MaxPooling2D(pool_size=(2, 2)),

        # Third conv layer
        Conv2D(256, kernel_size=(6, 6), strides=(1, 1), activation='relu', padding='same', kernel_regularizer=l2(0.002)),
        MaxPooling2D(pool_size=(2, 2)),

        # Fourth conv layer
        Conv2D(64, kernel_size=(3, 3), strides=(1, 1), activation='relu', padding='same', kernel_regularizer=l2(0.002)),
        MaxPooling2D(pool_size=(2, 2)),

        # Fifth conv layer
        Conv2D(32, kernel_size=(6, 6), strides=(1, 1), activation='relu', padding='same', kernel_regularizer=l2(0.002)),
        MaxPooling2D(pool_size=(2, 2)),

        # Flattening the layers
        Flatten(),

        # Dense layers
        Dense(512, activation='relu', kernel_regularizer=l2(0.002)),
        Dropout(0.5),
        Dense(256, activation='relu', kernel_regularizer=l2(0.002)),
        Dropout(0.5),
        Dense(num_classes, activation='softmax')
    ])

    model.compile(optimizer=Adam(), loss='categorical_crossentropy', metrics=['accuracy'])
    return model
```

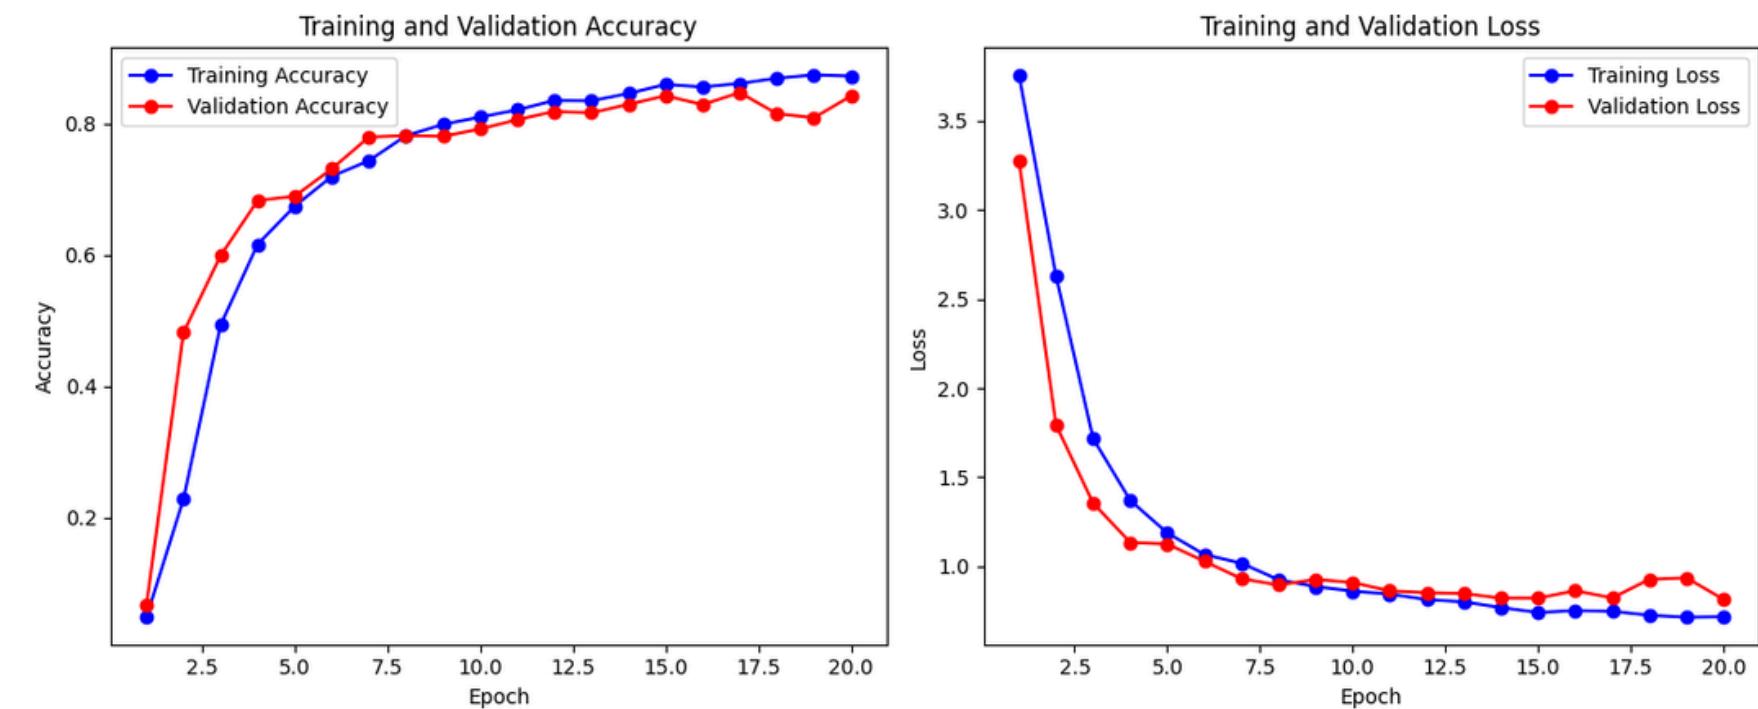
MODEL SUMMARY

Model: "sequential_2"		
Layer (type)	Output Shape	Param #
conv2d_10 (Conv2D)	(None, 32, 32, 64)	640
max_pooling2d_10 (MaxPooling2D)	(None, 16, 16, 64)	0
conv2d_11 (Conv2D)	(None, 16, 16, 32)	73760
max_pooling2d_11 (MaxPooling2D)	(None, 8, 8, 32)	0
conv2d_12 (Conv2D)	(None, 8, 8, 256)	295168
max_pooling2d_12 (MaxPooling2D)	(None, 4, 4, 256)	0
conv2d_13 (Conv2D)	(None, 4, 4, 64)	147520
max_pooling2d_13 (MaxPooling2D)	(None, 2, 2, 64)	0
conv2d_14 (Conv2D)	(None, 2, 2, 32)	73760
max_pooling2d_14 (MaxPooling2D)	(None, 1, 1, 32)	0
flatten_2 (Flatten)	(None, 32)	0
dense_6 (Dense)	(None, 512)	16896
dropout_4 (Dropout)	(None, 512)	0
dense_7 (Dense)	(None, 256)	131328
dropout_5 (Dropout)	(None, 256)	0
dense_8 (Dense)	(None, 28)	7196
<hr/>		
Total params: 746268 (2.85 MB)		
Trainable params: 746268 (2.85 MB)		
Non-trainable params: 0 (0.00 Byte)		

TRAINING VS VALIDATION

```
history = model.fit(  
    train_images, train_labels_one_hot,  
    epochs=20,  
    batch_size=32,  
    validation_split=0.2, # Using 20% of the data for validation  
    verbose=1  
)  
  
Epoch 1/20  
336/336 [=====] - 6s 10ms/step - loss: 3.7567 - accuracy: 0.0484 - val_loss: 3.2730 - val_accuracy: 0.0666  
Epoch 2/20  
336/336 [=====] - 3s 8ms/step - loss: 2.6273 - accuracy: 0.2286 - val_loss: 1.7924 - val_accuracy: 0.4821  
Epoch 3/20  
336/336 [=====] - 2s 7ms/step - loss: 1.7165 - accuracy: 0.4936 - val_loss: 1.3547 - val_accuracy: 0.5993  
Epoch 4/20  
336/336 [=====] - 2s 7ms/step - loss: 1.3704 - accuracy: 0.6164 - val_loss: 1.1346 - val_accuracy: 0.6830  
Epoch 5/20  
336/336 [=====] - 2s 7ms/step - loss: 1.1876 - accuracy: 0.6746 - val_loss: 1.1258 - val_accuracy: 0.6897  
Epoch 6/20  
336/336 [=====] - 2s 7ms/step - loss: 1.0667 - accuracy: 0.7199 - val_loss: 1.0285 - val_accuracy: 0.7314  
Epoch 7/20  
336/336 [=====] - 3s 9ms/step - loss: 1.0178 - accuracy: 0.7439 - val_loss: 0.9313 - val_accuracy: 0.7798  
Epoch 8/20  
336/336 [=====] - 3s 8ms/step - loss: 0.9253 - accuracy: 0.7818 - val_loss: 0.8957 - val_accuracy: 0.7820  
Epoch 9/20  
336/336 [=====] - 2s 7ms/step - loss: 0.8870 - accuracy: 0.7991 - val_loss: 0.9272 - val_accuracy: 0.7809  
Epoch 10/20  
336/336 [=====] - 2s 7ms/step - loss: 0.8614 - accuracy: 0.8102 - val_loss: 0.9096 - val_accuracy: 0.7920  
Epoch 11/20  
336/336 [=====] - 2s 7ms/step - loss: 0.8450 - accuracy: 0.8211 - val_loss: 0.8630 - val_accuracy: 0.8062  
Epoch 12/20  
336/336 [=====] - 2s 7ms/step - loss: 0.8140 - accuracy: 0.8355 - val_loss: 0.8528 - val_accuracy: 0.8188  
Epoch 13/20  
336/336 [=====] - 3s 9ms/step - loss: 0.8011 - accuracy: 0.8350 - val_loss: 0.8477 - val_accuracy: 0.8166  
Epoch 14/20  
336/336 [=====] - 3s 9ms/step - loss: 0.7694 - accuracy: 0.8461 - val_loss: 0.8225 - val_accuracy: 0.8296  
Epoch 15/20  
336/336 [=====] - 2s 7ms/step - loss: 0.7417 - accuracy: 0.8599 - val_loss: 0.8220 - val_accuracy: 0.8426  
Epoch 16/20  
336/336 [=====] - 2s 7ms/step - loss: 0.7524 - accuracy: 0.8560 - val_loss: 0.8640 - val_accuracy: 0.8292  
Epoch 17/20  
336/336 [=====] - 2s 7ms/step - loss: 0.7479 - accuracy: 0.8614 - val_loss: 0.8234 - val_accuracy: 0.8475  
Epoch 18/20  
336/336 [=====] - 2s 7ms/step - loss: 0.7263 - accuracy: 0.8692 - val_loss: 0.9274 - val_accuracy: 0.8155  
Epoch 19/20  
336/336 [=====] - 3s 8ms/step - loss: 0.7149 - accuracy: 0.8745 - val_loss: 0.9363 - val_accuracy: 0.8092  
Epoch 20/20  
336/336 [=====] - 3s 9ms/step - loss: 0.7179 - accuracy: 0.8727 - val_loss: 0.8153 - val_accuracy: 0.8423
```

training accuracy: 0.87
validation accuracy: 0.84

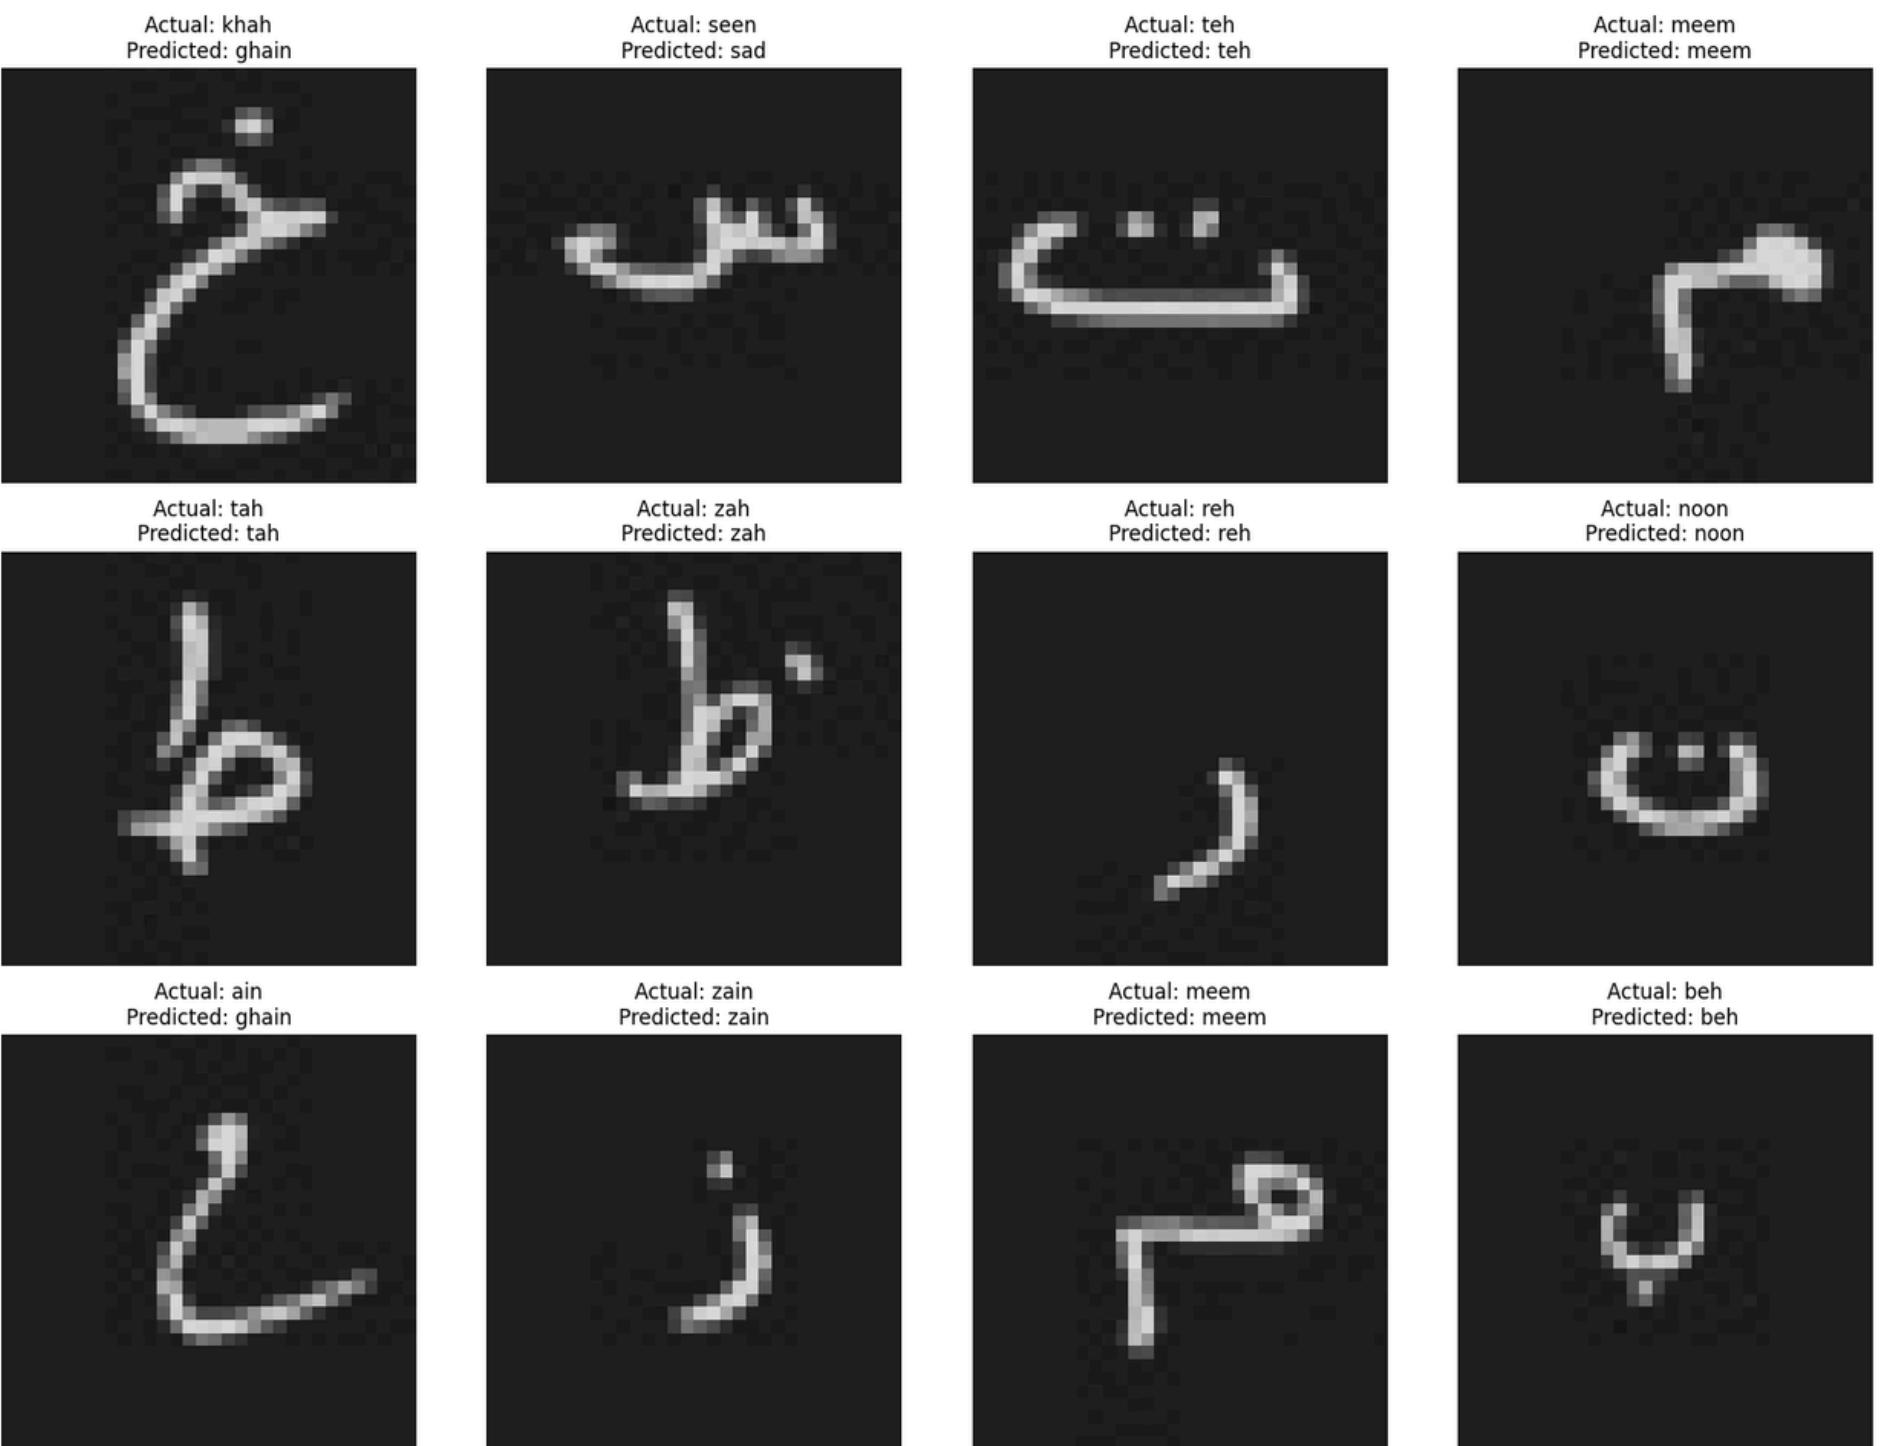


TESTING

```
# Evaluate the model on the test set
test_loss, test_accuracy = model.evaluate(test_images, test_labels_one_hot)
print(f"Test Loss: {test_loss}")
print(f"Test Accuracy: {test_accuracy}")

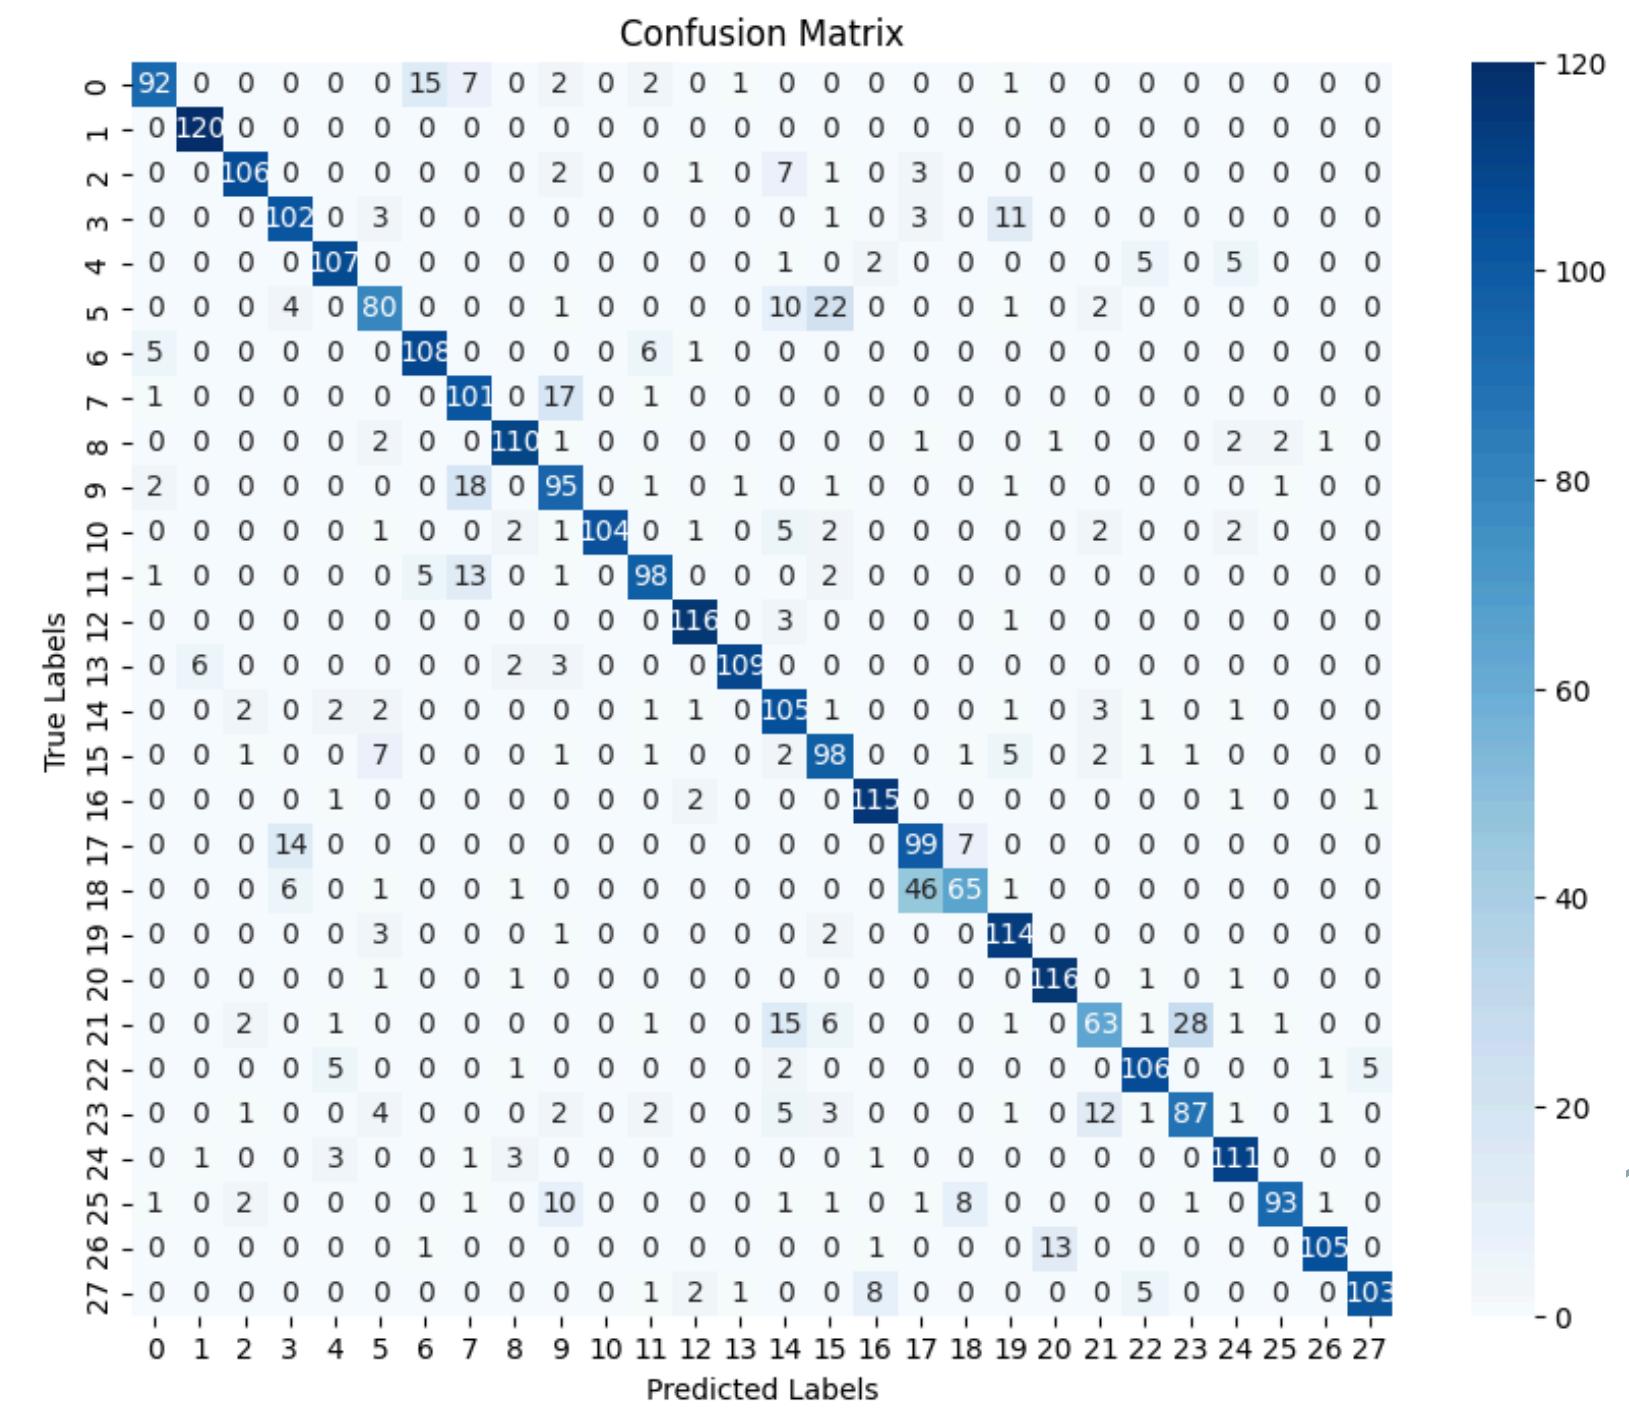
105/105 [=====] - 0s 4ms/step - loss: 0.8378 - accuracy: 0.8417
Test Loss: 0.8377987146377563
Test Accuracy: 0.8416666388511658
```

testing accuracy: 0.84



CLASSIFICATION REPORT & CONFUSION MATRIX

	precision	recall	f1-score	support
0	0.90	0.77	0.83	120
1	0.94	1.00	0.97	120
2	0.93	0.88	0.91	120
3	0.81	0.85	0.83	120
4	0.90	0.89	0.90	120
5	0.77	0.67	0.71	120
6	0.84	0.90	0.87	120
7	0.72	0.84	0.77	120
8	0.92	0.92	0.92	120
9	0.69	0.79	0.74	120
10	1.00	0.87	0.93	120
11	0.86	0.82	0.84	120
12	0.94	0.97	0.95	120
13	0.97	0.91	0.94	120
14	0.67	0.88	0.76	120
15	0.70	0.82	0.75	120
16	0.91	0.96	0.93	120
17	0.65	0.82	0.73	120
18	0.80	0.54	0.65	120
19	0.83	0.95	0.88	120
20	0.89	0.97	0.93	120
21	0.75	0.53	0.62	120
22	0.88	0.88	0.88	120
23	0.74	0.72	0.73	120
24	0.89	0.93	0.91	120
25	0.96	0.78	0.86	120
26	0.96	0.88	0.92	120
27	0.94	0.86	0.90	120
accuracy			0.84	3360
macro avg	0.85	0.84	0.84	3360
weighted avg	0.85	0.84	0.84	3360



CHANGE 2 TO THE MODEL

changing convolution size, kernel size, L2 regularization
&
Early Stopping

IMPLEMENTATION

```
def create_advanced_cnn_model(input_shape, num_classes):
    model = Sequential([
        # First conv layer
        Conv2D(128, kernel_size=(5, 5), strides=(1, 1), activation='relu', input_shape=input_shape, padding='same', kernel_regularizer=l2(0.01)),
        MaxPooling2D(pool_size=(2, 2)),

        # Second conv layer
        Conv2D(256, kernel_size=(5, 5), strides=(1, 1), activation='relu', padding='same', kernel_regularizer=l2(0.01)),
        MaxPooling2D(pool_size=(2, 2)),

        # Third conv layer
        Conv2D(512, kernel_size=(3, 3), strides=(1, 1), activation='relu', padding='same', kernel_regularizer=l2(0.01)),
        MaxPooling2D(pool_size=(2, 2)),

        # Fourth conv layer
        Conv2D(512, kernel_size=(3, 3), strides=(1, 1), activation='relu', padding='same', kernel_regularizer=l2(0.01)),
        MaxPooling2D(pool_size=(2, 2)),

        # Fifth conv layer
        Conv2D(512, kernel_size=(3, 3), strides=(1, 1), activation='relu', padding='same', kernel_regularizer=l2(0.01)),
        MaxPooling2D(pool_size=(2, 2)),

        # Flattening the layers
        Flatten(),

        # Dense layers
        Dense(1024, activation='relu', kernel_regularizer=l2(0.01)),
        Dropout(0.5),
        Dense(512, activation='relu', kernel_regularizer=l2(0.01)),
        Dropout(0.5),
        Dense(num_classes, activation='softmax')
    ])

    model.compile(optimizer=Adam(), loss='categorical_crossentropy', metrics=['accuracy'])

    # Adding Early Stopping callback
    early_stopping = EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True)

    return model, early_stopping
```

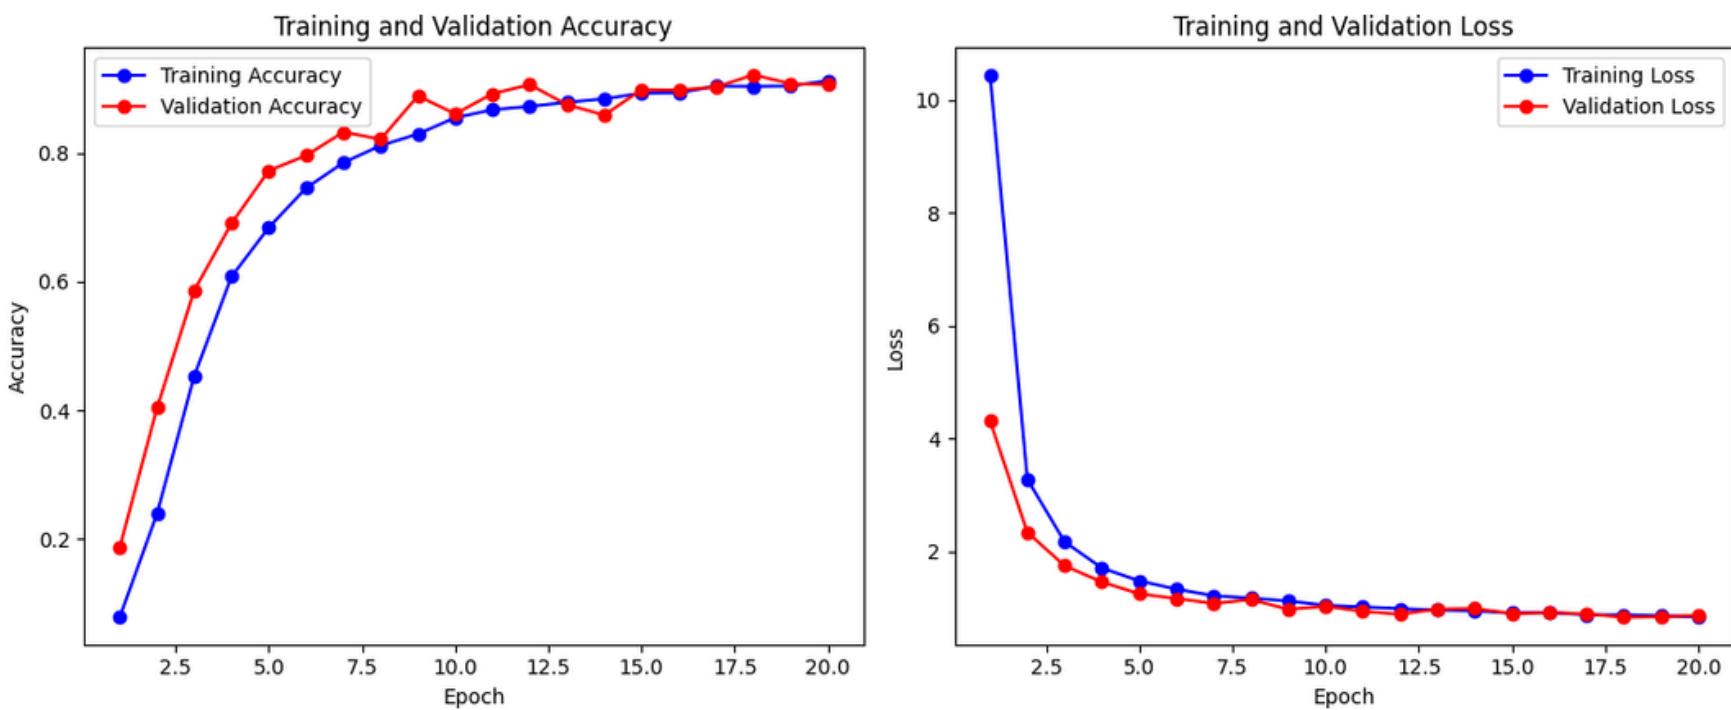
MODEL SUMMARY

Model: "sequential_5"		
Layer (type)	Output Shape	Param #
conv2d_25 (Conv2D)	(None, 32, 32, 128)	3328
max_pooling2d_25 (MaxPooling2D)	(None, 16, 16, 128)	0
conv2d_26 (Conv2D)	(None, 16, 16, 256)	819456
max_pooling2d_26 (MaxPooling2D)	(None, 8, 8, 256)	0
conv2d_27 (Conv2D)	(None, 8, 8, 512)	1180160
max_pooling2d_27 (MaxPooling2D)	(None, 4, 4, 512)	0
conv2d_28 (Conv2D)	(None, 4, 4, 512)	2359808
max_pooling2d_28 (MaxPooling2D)	(None, 2, 2, 512)	0
conv2d_29 (Conv2D)	(None, 2, 2, 512)	2359808
max_pooling2d_29 (MaxPooling2D)	(None, 1, 1, 512)	0
flatten_5 (Flatten)	(None, 512)	0
dense_15 (Dense)	(None, 1024)	525312
dropout_10 (Dropout)	(None, 1024)	0
dense_16 (Dense)	(None, 512)	524800
dropout_11 (Dropout)	(None, 512)	0
dense_17 (Dense)	(None, 28)	14364
<hr/>		
Total params: 7787036 (29.71 MB)		
Trainable params: 7787036 (29.71 MB)		
Non-trainable params: 0 (0.00 Byte)		

TRAINING VS VALIDATION

```
history = model.fit(  
    train_images, train_labels_one_hot,  
    epochs=20,  
    batch_size=32,  
    validation_split=0.2, # Using 20% of the data for validation  
    verbose=1  
)  
  
Epoch 1/20  
336/336 [=====] - 11s 22ms/step - loss: 10.4446 - accuracy: 0.0789 - val_loss: 4.3122 - val_accuracy: 0.1871  
Epoch 2/20  
336/336 [=====] - 7s 22ms/step - loss: 3.2791 - accuracy: 0.2398 - val_loss: 2.3394 - val_accuracy: 0.4040  
Epoch 3/20  
336/336 [=====] - 7s 22ms/step - loss: 2.1712 - accuracy: 0.4540 - val_loss: 1.7466 - val_accuracy: 0.5867  
Epoch 4/20  
336/336 [=====] - 8s 22ms/step - loss: 1.7061 - accuracy: 0.6082 - val_loss: 1.4585 - val_accuracy: 0.6908  
Epoch 5/20  
336/336 [=====] - 7s 22ms/step - loss: 1.4795 - accuracy: 0.6844 - val_loss: 1.2502 - val_accuracy: 0.7727  
Epoch 6/20  
336/336 [=====] - 7s 21ms/step - loss: 1.3325 - accuracy: 0.7456 - val_loss: 1.1662 - val_accuracy: 0.7961  
Epoch 7/20  
336/336 [=====] - 7s 21ms/step - loss: 1.2157 - accuracy: 0.7851 - val_loss: 1.0806 - val_accuracy: 0.8326  
Epoch 8/20  
336/336 [=====] - 7s 21ms/step - loss: 1.1742 - accuracy: 0.8116 - val_loss: 1.1469 - val_accuracy: 0.8214  
Epoch 9/20  
336/336 [=====] - 7s 21ms/step - loss: 1.1223 - accuracy: 0.8296 - val_loss: 0.9756 - val_accuracy: 0.8888  
Epoch 10/20  
336/336 [=====] - 7s 20ms/step - loss: 1.0459 - accuracy: 0.8549 - val_loss: 1.0275 - val_accuracy: 0.8609  
Epoch 11/20  
336/336 [=====] - 7s 22ms/step - loss: 1.0167 - accuracy: 0.8673 - val_loss: 0.9394 - val_accuracy: 0.8921  
Epoch 12/20  
336/336 [=====] - 7s 20ms/step - loss: 0.9865 - accuracy: 0.8725 - val_loss: 0.8881 - val_accuracy: 0.9062  
Epoch 13/20  
336/336 [=====] - 7s 22ms/step - loss: 0.9625 - accuracy: 0.8788 - val_loss: 0.9801 - val_accuracy: 0.8750  
Epoch 14/20  
336/336 [=====] - 7s 20ms/step - loss: 0.9483 - accuracy: 0.8843 - val_loss: 0.9882 - val_accuracy: 0.8590  
Epoch 15/20  
336/336 [=====] - 7s 22ms/step - loss: 0.9087 - accuracy: 0.8929 - val_loss: 0.9007 - val_accuracy: 0.8984  
Epoch 16/20  
336/336 [=====] - 7s 21ms/step - loss: 0.9185 - accuracy: 0.8935 - val_loss: 0.9169 - val_accuracy: 0.8977  
Epoch 17/20  
336/336 [=====] - 7s 21ms/step - loss: 0.8769 - accuracy: 0.9040 - val_loss: 0.8936 - val_accuracy: 0.9025  
Epoch 18/20  
336/336 [=====] - 7s 21ms/step - loss: 0.8775 - accuracy: 0.9036 - val_loss: 0.8343 - val_accuracy: 0.9215  
Epoch 19/20  
336/336 [=====] - 7s 21ms/step - loss: 0.8610 - accuracy: 0.9042 - val_loss: 0.8435 - val_accuracy: 0.9077  
Epoch 20/20  
336/336 [=====] - 7s 21ms/step - loss: 0.8393 - accuracy: 0.9124 - val_loss: 0.8673 - val_accuracy: 0.9062
```

training accuracy: 0.91
validation accuracy: 0.90

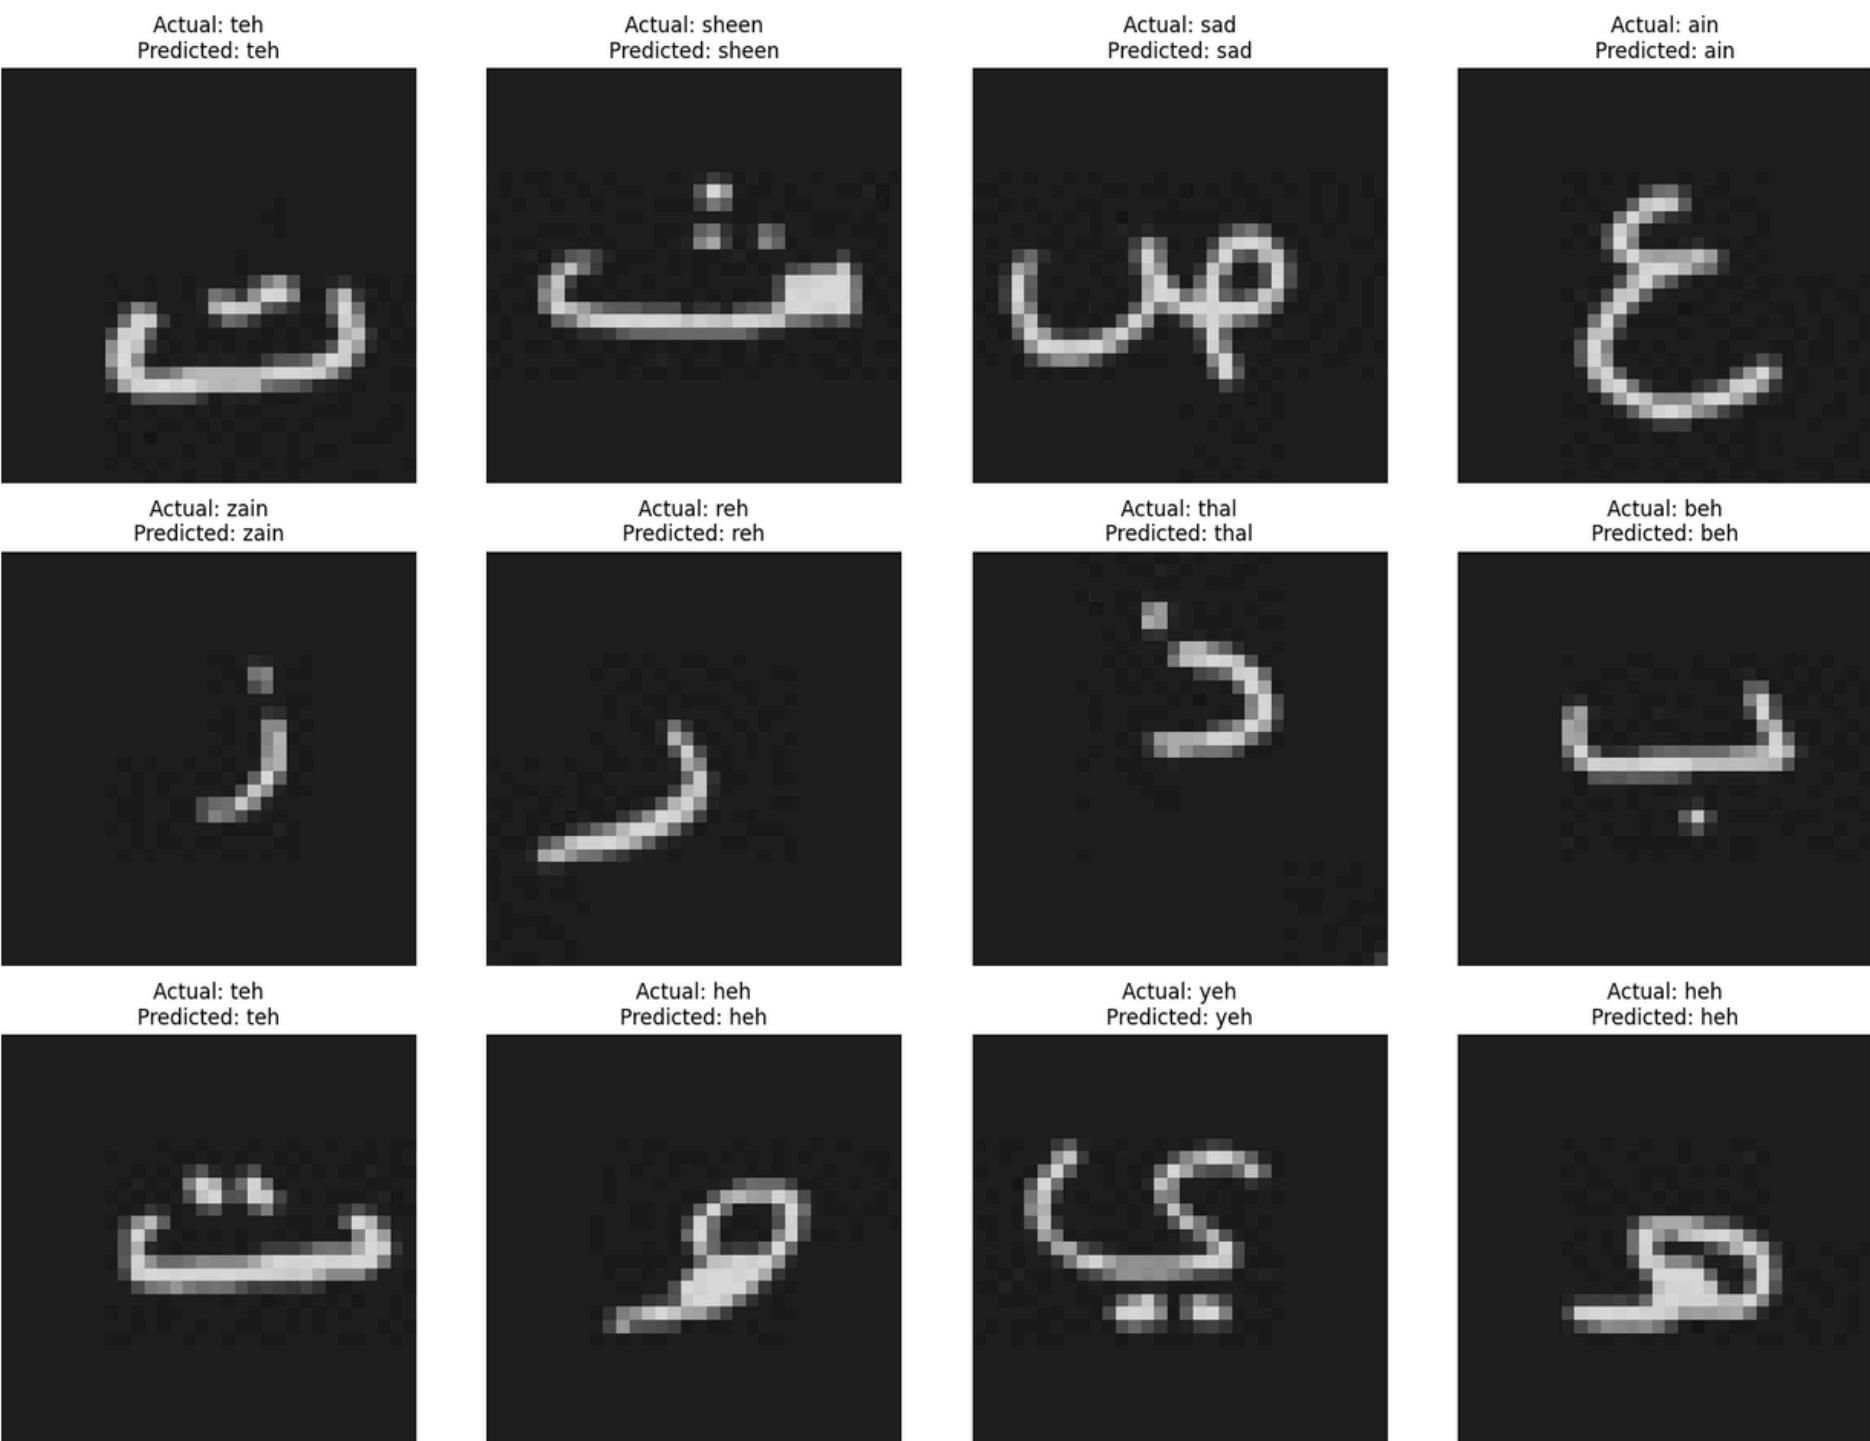


TESTING

```
# Evaluate the model on the test set
test_loss, test_accuracy = model.evaluate(test_images, test_labels_one_hot)
print(f"Test Loss: {test_loss}")
print(f"Test Accuracy: {test_accuracy}")

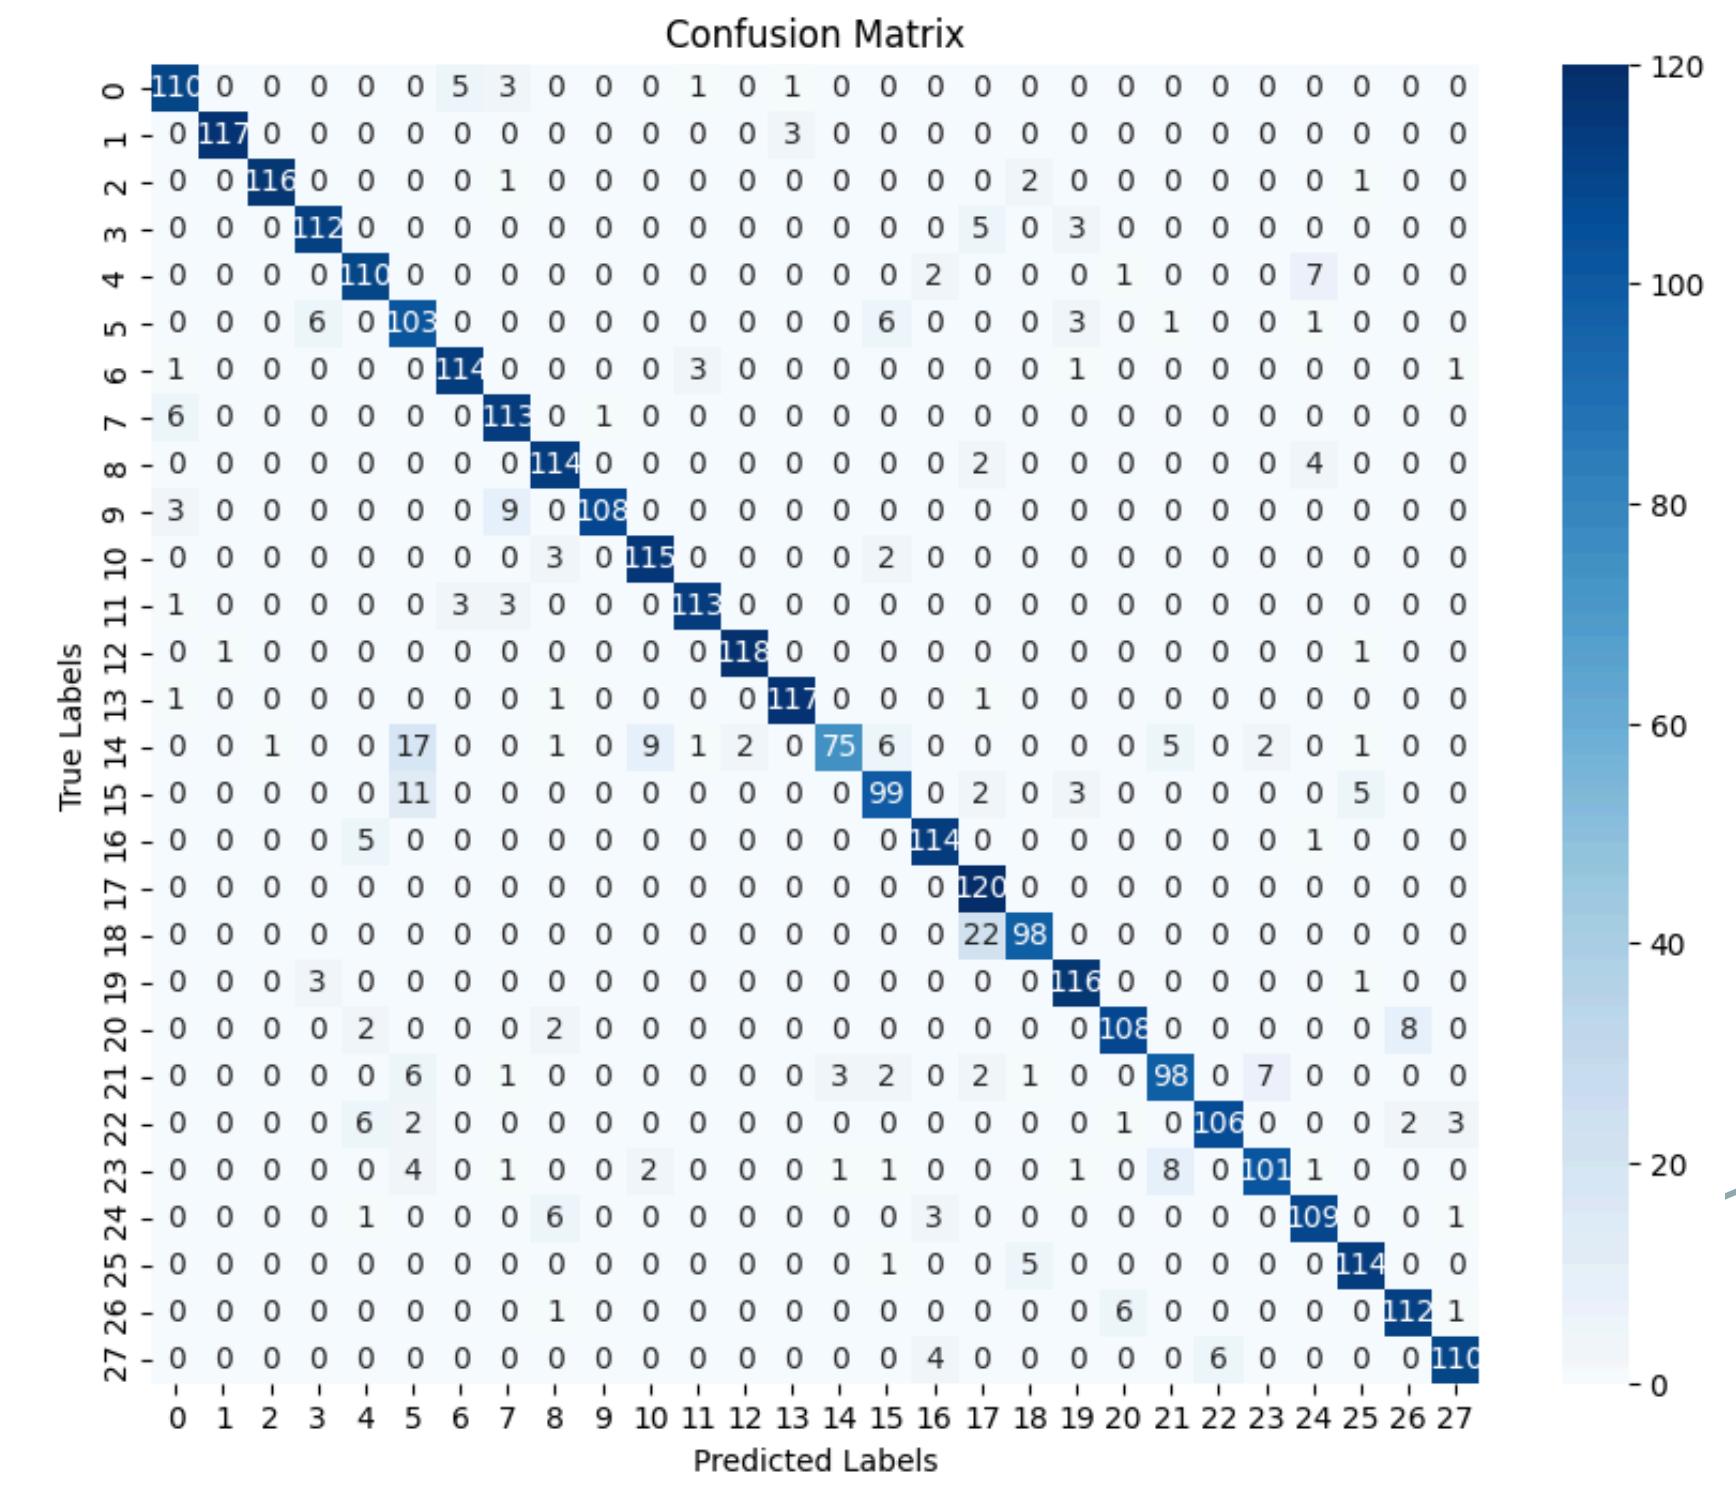
105/105 [=====] - 1s 7ms/step - loss: 0.8369 - accuracy: 0.9107
Test Loss: 0.8368542194366455
Test Accuracy: 0.9107142686843872
```

testing accuracy: 0.91



CLASSIFICATION REPORT & CONFUSION MATRIX

105/105 [=====] - 1s 5ms/step		precision	recall	f1-score	support
0		0.90	0.92	0.91	120
1		0.99	0.97	0.98	120
2		0.99	0.97	0.98	120
3		0.93	0.93	0.93	120
4		0.89	0.92	0.90	120
5		0.72	0.86	0.78	120
6		0.93	0.95	0.94	120
7		0.86	0.94	0.90	120
8		0.89	0.95	0.92	120
9		0.99	0.90	0.94	120
10		0.91	0.96	0.93	120
11		0.96	0.94	0.95	120
12		0.98	0.98	0.98	120
13		0.97	0.97	0.97	120
14		0.95	0.62	0.75	120
15		0.85	0.82	0.84	120
16		0.93	0.95	0.94	120
17		0.78	1.00	0.88	120
18		0.92	0.82	0.87	120
19		0.91	0.97	0.94	120
20		0.93	0.90	0.92	120
21		0.88	0.82	0.84	120
22		0.95	0.88	0.91	120
23		0.92	0.84	0.88	120
24		0.89	0.91	0.90	120
25		0.93	0.95	0.94	120
26		0.92	0.93	0.93	120
27		0.95	0.92	0.93	120
accuracy				0.91	3360
macro avg		0.91	0.91	0.91	3360
weighted avg		0.91	0.91	0.91	3360



CHANGE 3 TO THE MODEL

The use of Keras Tuner library

IMPLEMENTATION

```
def create_advanced_cnn_model(hp):
    model = Sequential()

    # Define hyperparameters
    hp_units = hp.Int('units', min_value=32, max_value=1024, step=32)
    hp_dropout = hp.Float('dropout', min_value=0.1, max_value=0.5, step=0.1)
    hp_learning_rate = hp.Choice('learning_rate', values=[1e-2, 1e-3, 1e-4])

    # First conv layer
    model.add(Conv2D(64, kernel_size=(3, 3), strides=(1, 1), activation='relu', input_shape=input_shape, padding='same', kernel_regularizer=l2(0.001)))
    model.add(MaxPooling2D(pool_size=(2, 2)))

    # Second conv layer
    model.add(Conv2D(128, kernel_size=(3, 3), strides=(1, 1), activation='relu', padding='same', kernel_regularizer=l2(0.001)))
    model.add(MaxPooling2D(pool_size=(2, 2)))

    # Third conv layer
    model.add(Conv2D(256, kernel_size=(3, 3), strides=(1, 1), activation='relu', padding='same', kernel_regularizer=l2(0.001)))
    model.add(MaxPooling2D(pool_size=(2, 2)))

    # Fourth conv layer
    model.add(Conv2D(512, kernel_size=(3, 3), strides=(1, 1), activation='relu', padding='same', kernel_regularizer=l2(0.001)))
    model.add(MaxPooling2D(pool_size=(2, 2)))

    # Fifth conv layer
    model.add(Conv2D(512, kernel_size=(3, 3), strides=(1, 1), activation='relu', padding='same', kernel_regularizer=l2(0.001)))
    model.add(MaxPooling2D(pool_size=(2, 2)))

    # Flattening the layers
    model.add(Flatten())

    # Dense layers with hyperparameters
    model.add(Dense(units=hp_units, activation='relu', kernel_regularizer=l2(0.001)))
    model.add(Dropout(hp_dropout))
    model.add(Dense(units=hp_units//2, activation='relu', kernel_regularizer=l2(0.001)))
    model.add(Dropout(hp_dropout))

    # Output layer
    model.add(Dense(num_classes, activation='softmax'))

    # Compile the model with the defined learning rate
    model.compile(optimizer=Adam(learning_rate=hp_learning_rate), loss='categorical_crossentropy', metrics=['accuracy'])

    return model
```

```
class CNNHyperModel(HyperModel):
    def __init__(self, input_shape, num_classes):
        self.input_shape = input_shape
        self.num_classes = num_classes

    def build(self, hp):
        return create_advanced_cnn_model(hp)
```

```
hypermodel = CNNHyperModel(input_shape, num_classes)
tuner = Hyperband(
    hypermodel,
    objective='val_accuracy',
    max_epochs=10,
    directory='tuner_dir',
    project_name='cnn_hyperband'
)
```

```
tuner.search(x_train, y_train, epochs=10, validation_data=(x_val, y_val))

Trial 30 Complete [00h 01m 25s]
val_accuracy: 0.9367559552192688

Best val_accuracy So Far: 0.9419642686843872
Total elapsed time: 00h 12m 37s
```

```
# Retrieve the best hyperparameters
best_hyperparameters = tuner.get_best_hyperparameters(1)[0]
best_model = tuner.hypermodel.build(best_hyperparameters)
print('the best hyperparameter is: ', best_hyperparameters)
print('the best model is: ', best_model)

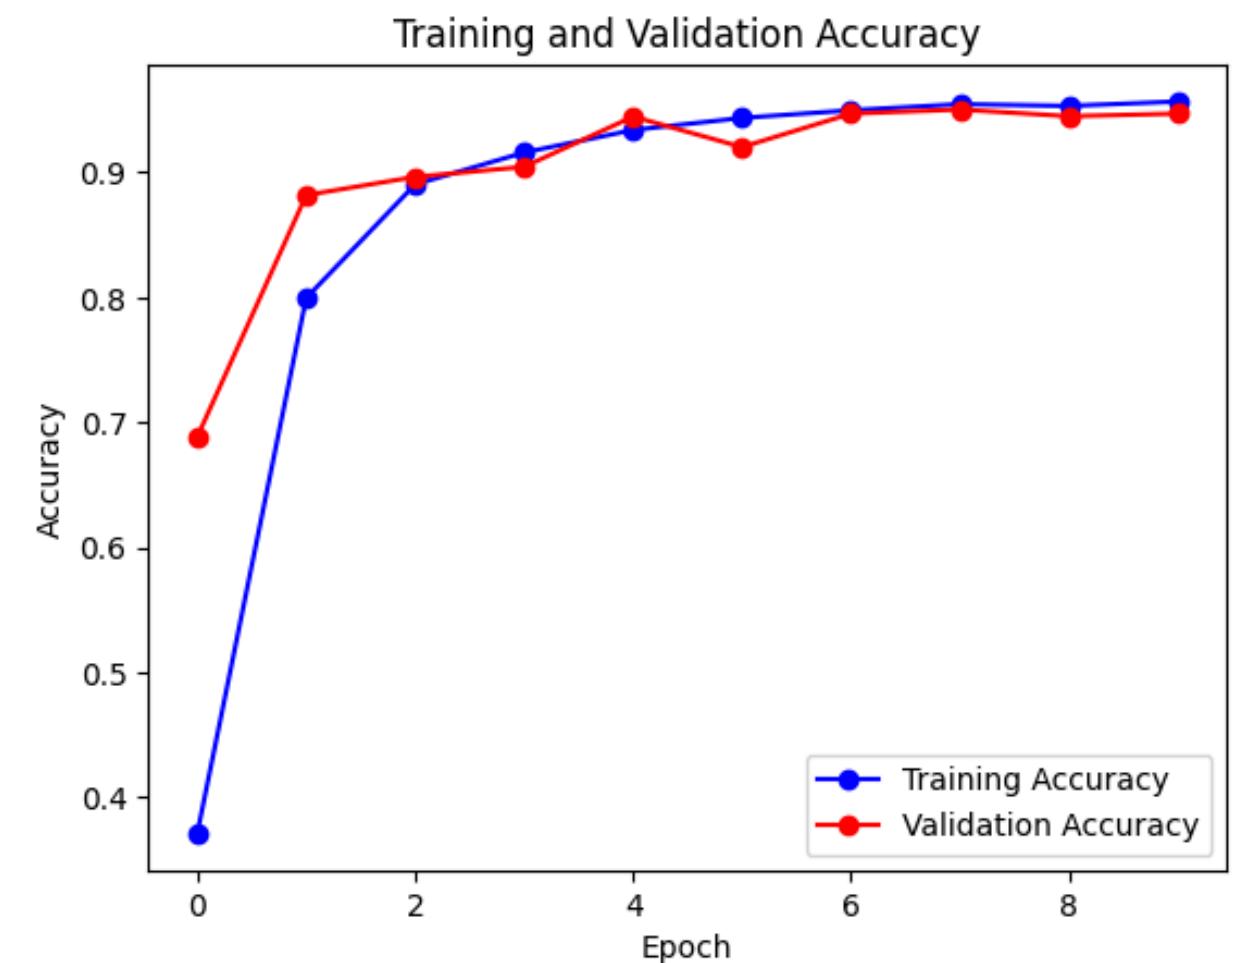
the best hyperparameter is: <keras_tuner.src.engine.hyperparameters.HyperParameters object at 0x7e7a807dec50>
the best model is: <keras.src.engine.sequential.Sequential object at 0x7e7a8851be50>
```

TRAINING VS VALIDATION

```
# Train the best model on the full training dataset
history = best_model.fit(train_images, train_labels_one_hot, epochs=10, validation_data=(test_images, test_labels_one_hot))

Epoch 1/10
420/420 [=====] - 7s 12ms/step - loss: 3.1701 - accuracy: 0.3712 - val_loss: 1.6809 - val_accuracy: 0.6884
Epoch 2/10
420/420 [=====] - 5s 12ms/step - loss: 1.2753 - accuracy: 0.7996 - val_loss: 0.9502 - val_accuracy: 0.8818
Epoch 3/10
420/420 [=====] - 5s 12ms/step - loss: 0.8290 - accuracy: 0.8904 - val_loss: 0.7531 - val_accuracy: 0.8964
Epoch 4/10
420/420 [=====] - 5s 11ms/step - loss: 0.6553 - accuracy: 0.9159 - val_loss: 0.6798 - val_accuracy: 0.9045
Epoch 5/10
420/420 [=====] - 5s 12ms/step - loss: 0.5574 - accuracy: 0.9339 - val_loss: 0.5198 - val_accuracy: 0.9446
Epoch 6/10
420/420 [=====] - 5s 12ms/step - loss: 0.4935 - accuracy: 0.9436 - val_loss: 0.5540 - val_accuracy: 0.9202
Epoch 7/10
420/420 [=====] - 5s 11ms/step - loss: 0.4576 - accuracy: 0.9497 - val_loss: 0.4662 - val_accuracy: 0.9470
Epoch 8/10
420/420 [=====] - 5s 11ms/step - loss: 0.4294 - accuracy: 0.9546 - val_loss: 0.4600 - val_accuracy: 0.9503
Epoch 9/10
420/420 [=====] - 5s 13ms/step - loss: 0.4293 - accuracy: 0.9532 - val_loss: 0.4672 - val_accuracy: 0.9449
Epoch 10/10
420/420 [=====] - 5s 11ms/step - loss: 0.4094 - accuracy: 0.9567 - val_loss: 0.4414 - val_accuracy: 0.9470
```

training accuracy: 0.95
validation accuracy: 0.94

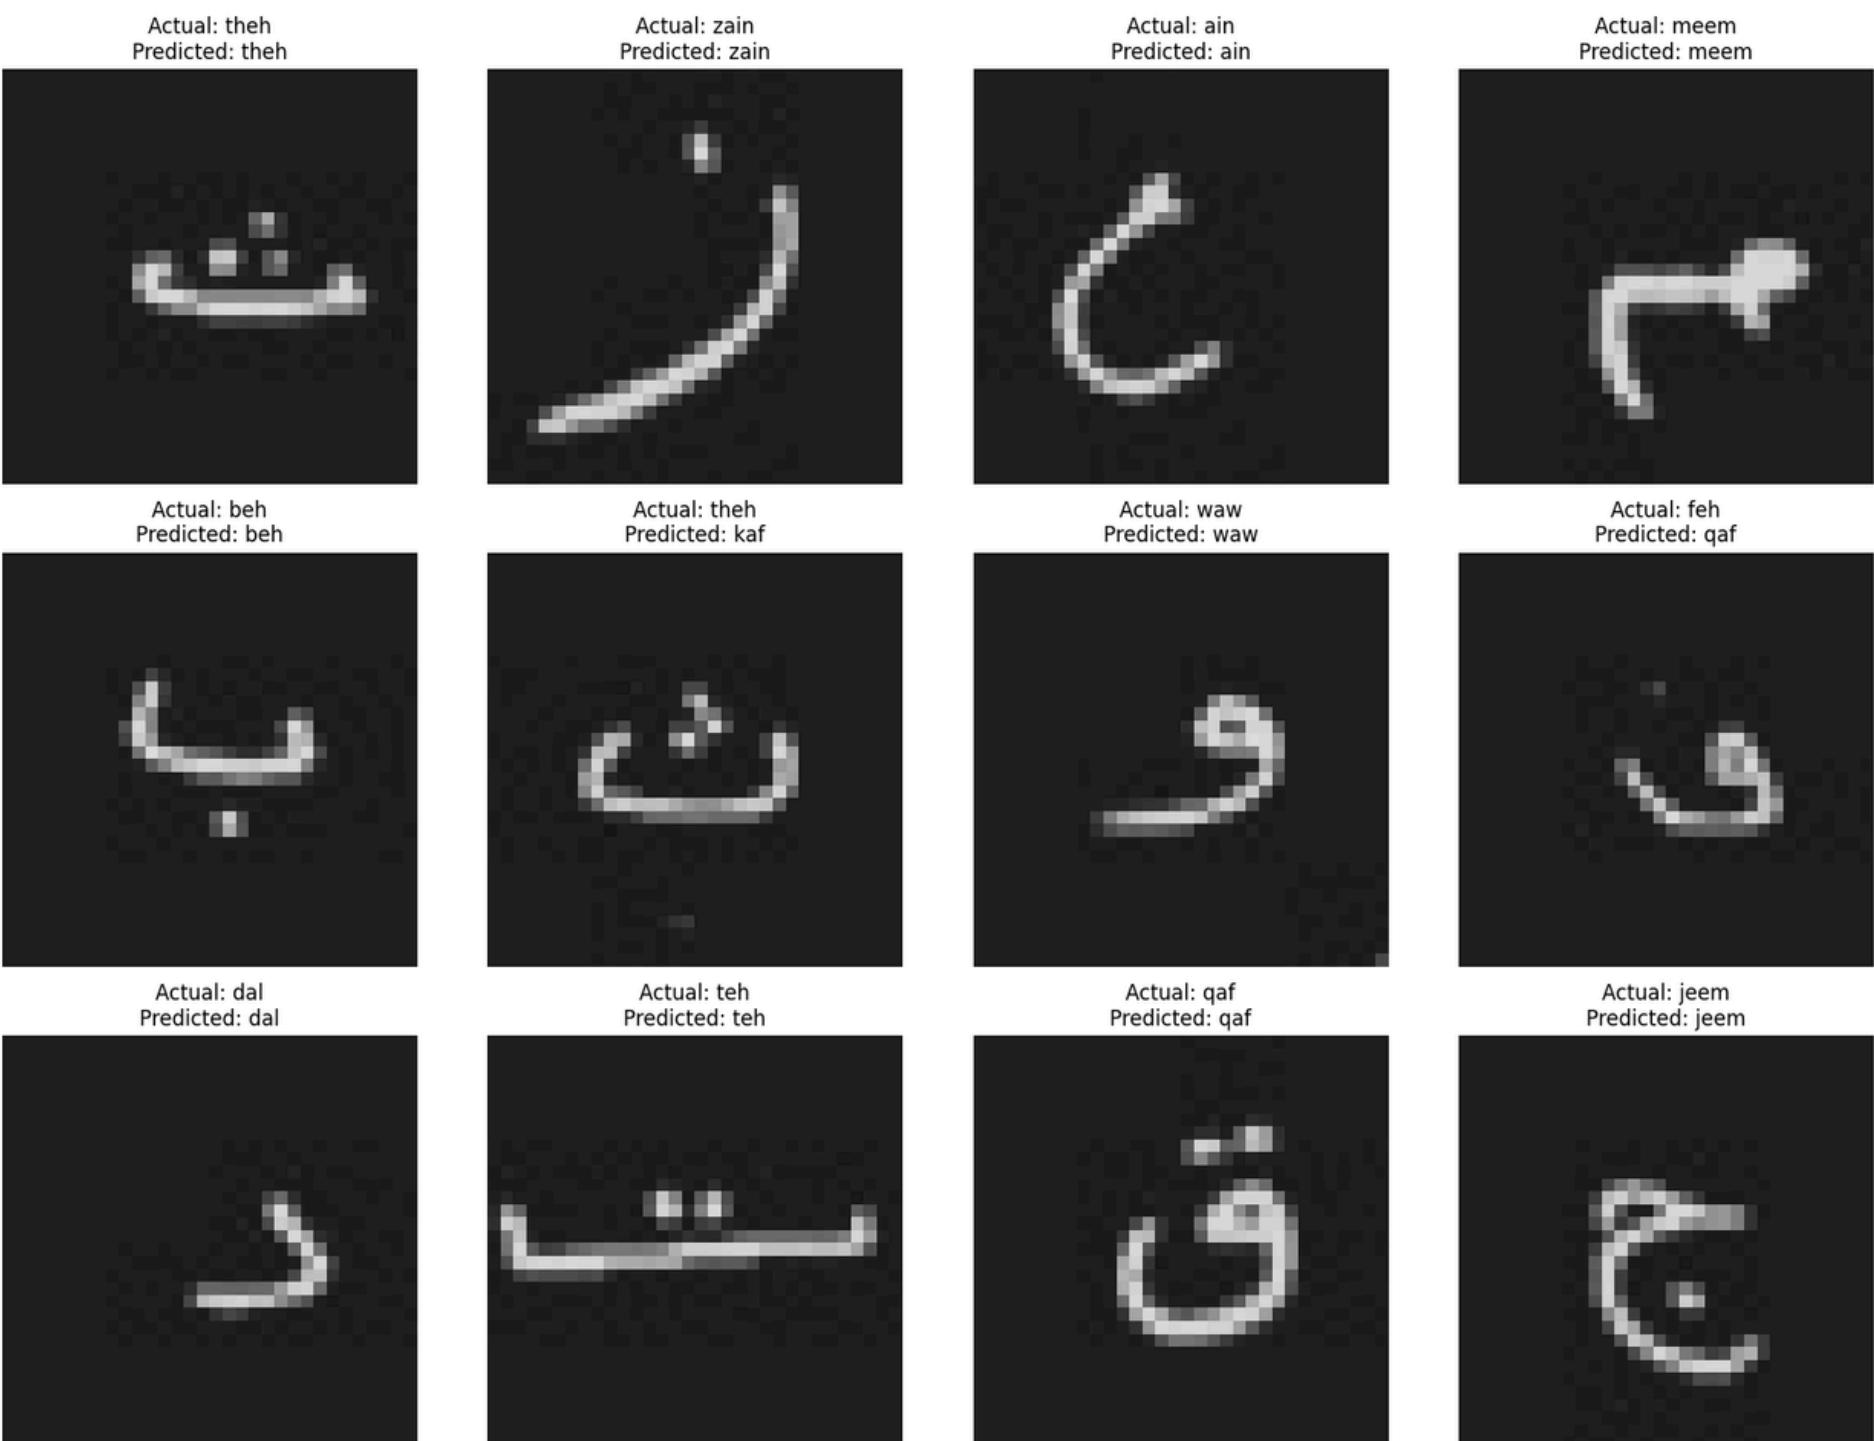


TESTING

```
# Evaluate the best model on the test dataset
test_loss, test_accuracy = best_model.evaluate(test_images, test_labels_one_hot)
print("Test Loss:", test_loss)
print("Test Accuracy:", test_accuracy)

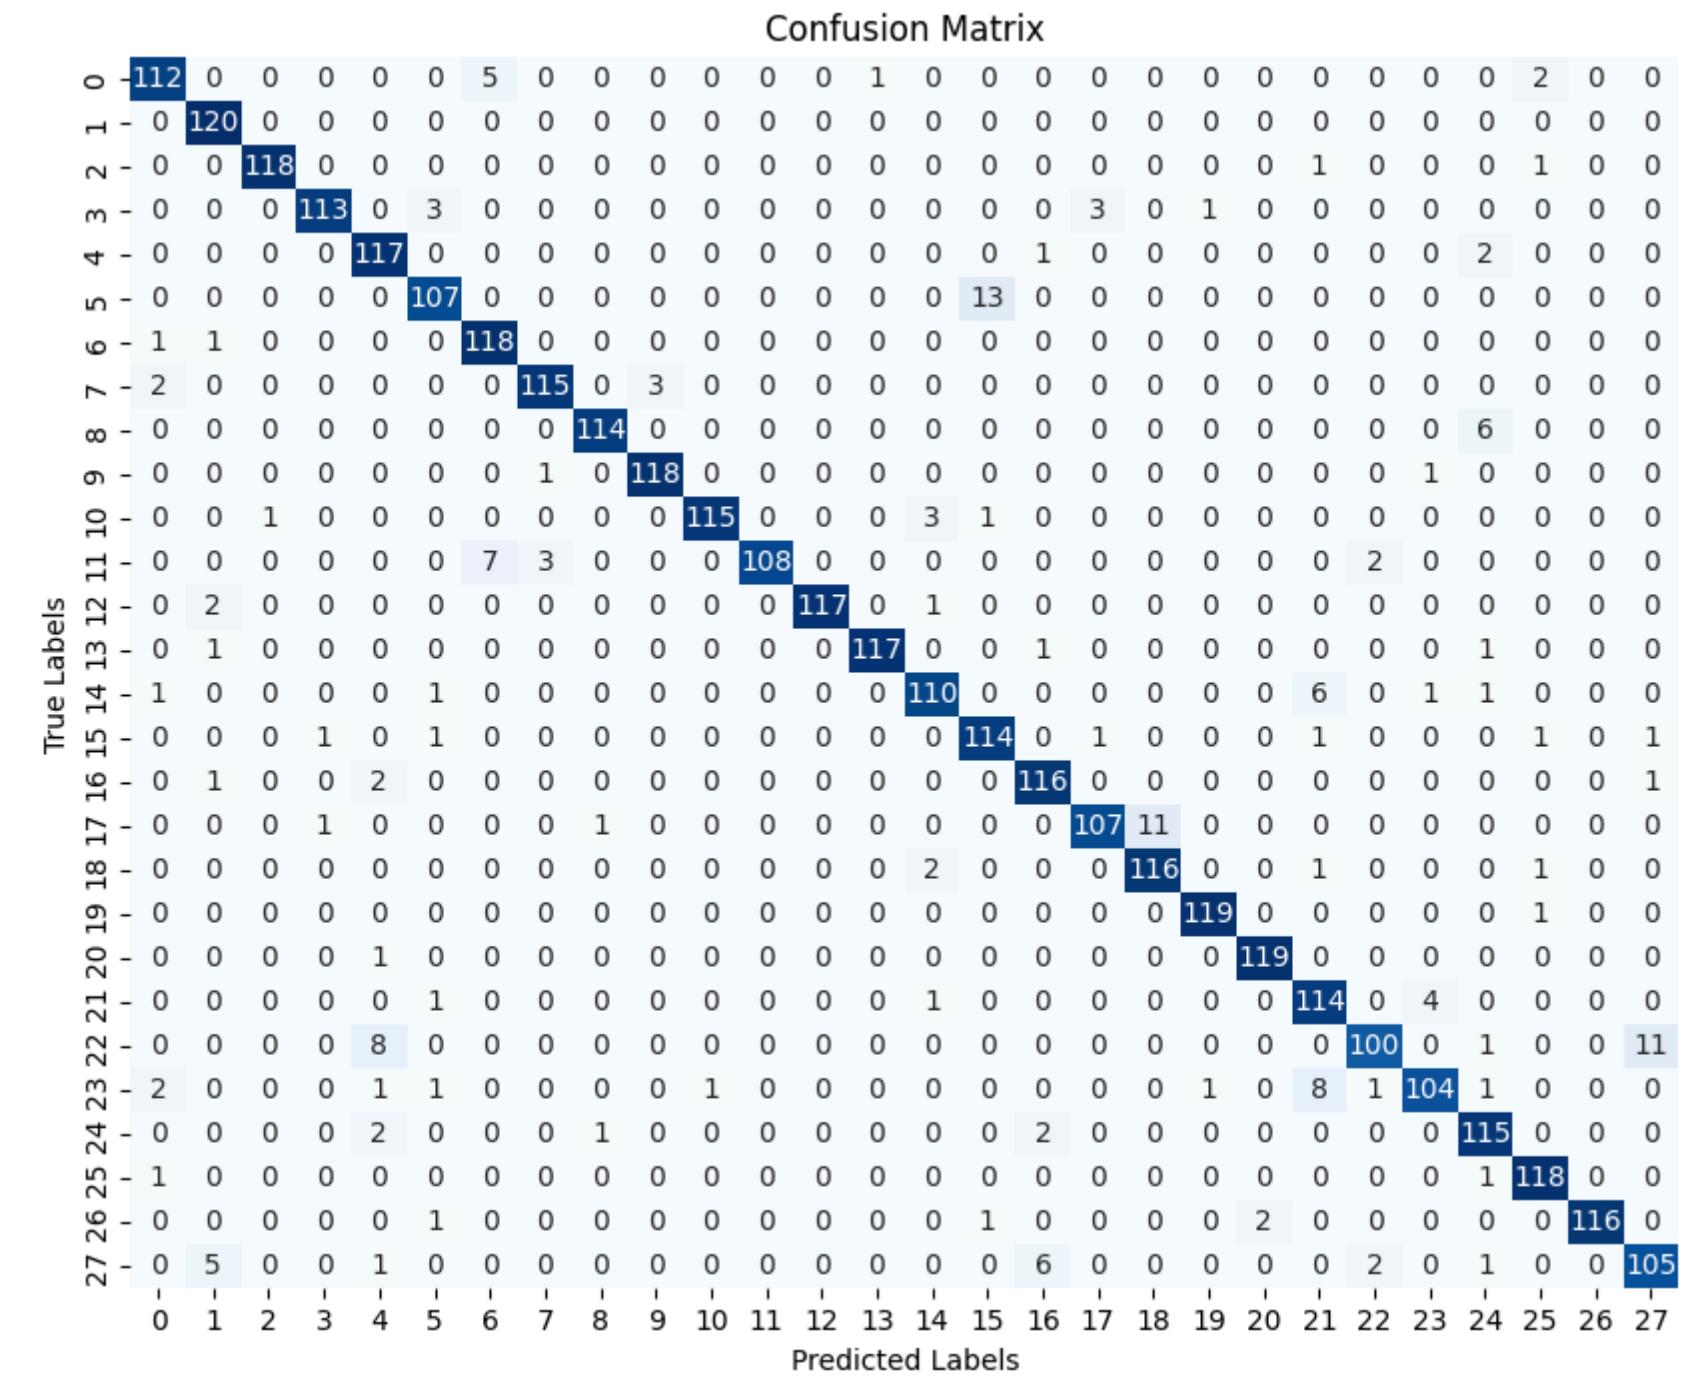
105/105 [=====] - 0s 4ms/step - loss: 0.4414 - accuracy: 0.9470
Test Loss: 0.4414333403110504
Test Accuracy: 0.9470238089561462
```

testing accuarcy: 0.94



CLASSIFICATION REPORT & CONFUSION MATRIX

105/105 [=====] - 0s 3ms/step				
	precision	recall	f1-score	support
0	0.94	0.93	0.94	120
1	0.92	1.00	0.96	120
2	0.99	0.98	0.99	120
3	0.98	0.94	0.96	120
4	0.89	0.97	0.93	120
5	0.93	0.89	0.91	120
6	0.91	0.98	0.94	120
7	0.97	0.96	0.96	120
8	0.98	0.95	0.97	120
9	0.98	0.98	0.98	120
10	0.99	0.96	0.97	120
11	1.00	0.90	0.95	120
12	1.00	0.97	0.99	120
13	0.99	0.97	0.98	120
14	0.94	0.92	0.93	120
15	0.88	0.95	0.92	120
16	0.92	0.97	0.94	120
17	0.96	0.89	0.93	120
18	0.91	0.97	0.94	120
19	0.98	0.99	0.99	120
20	0.98	0.99	0.99	120
21	0.87	0.95	0.91	120
22	0.95	0.83	0.89	120
23	0.95	0.87	0.90	120
24	0.89	0.96	0.92	120
25	0.95	0.98	0.97	120
26	1.00	0.97	0.98	120
27	0.89	0.88	0.88	120
accuracy			0.95	3360
macro avg		0.95	0.95	3360
weighted avg		0.95	0.95	3360



CONCLUSION

Model	Training	Validation	Testing
CNN model	0.97	0.92	0.93
CNN with changes 1	0.87	0.84	0.84
CNN with changes 2	0.91	0.90	0.91
CNN with changes 3	0.95	0.94	0.94

CNN Model: shows high accuracy across training, validation, and testing with scores of 0.97, 0.92, and 0.93, respectively. This model demonstrates a strong ability to generalize from training data to unseen validation and testing data, although there's a slight overfitting as indicated by the difference between training and validation/testing accuracy.

CNN Model change 1: has the lowest performance among the four, with a training accuracy of 0.87 and both validation and testing accuracies at 0.84. This model appears to be more consistent between training and external data results but has lower overall accuracy.

CNN Model change 2: shows good performance with training, validation, and testing accuracies of 0.91, 0.90, and 0.91, respectively. The consistency across these metrics suggests that the model generalizes well without significant overfitting.

CNN Model change 3: has the **highest performance** with a training accuracy of 0.95 and validation and testing accuracies of 0.94. It presents a very high level of generalization capability and minimal overfitting.

THANK YOU